



# Highlights/Overview

- 
- intro to AI/ML/DL/NNs
  - Hidden layers/Initialization/Neurons per layer
  - cost function/gradient descent/learning rate
  - Dropout rate
  - Activation function
  - Linear Regression
  - What are CNNs
  - Filters/ReLU/MaxPooling
  - Keras and CNNs
  - TensorFlow 1.x

# The Data/AI Landscape




# Use Cases for Deep Learning


- 
- computer vision
  - speech recognition
  - image processing
  - bioinformatics
  - social network filtering
  - drug design
  - Recommendation systems
  - Bioinformatics
  - Mobile Advertising
  - Many others
- 

# NN 3 Hidden Layers: Classifier

## Neural Networks



# NN: 2 Hidden Layers (Regression)





# Titanic Dataset (portion)


	A	B	C	D	E	F	G	H	I	J	K	L
1	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
2	1	0	3	Braund, Mr	male	22	1	0	A/5 21171	7.25		S
3	2	1	1	Cumings, Mr	female	38	1	0	PC 17599	71.2833	C85	C
4	3	1	3	Heikkinen, Mrs	female	26	0	0	STON/O2. 3101282	7.925		S
5	4	1	1	Futrelle, Mrs	female	35	1	0	113803	53.1	C123	S
6	5	0	3	Allen, Mr.	male	35	0	0	373450	8.05		S
7	6	0	3	Moran, Mr.	male		0	0	330877	8.4583		Q
8	7	0	1	McCarthy, Mr	male	54	0	0	17463	51.8625	E46	S
9	8	0	3	Palsson, Mr	male	2	3	1	349909	21.075		S
10	9	1	3	Johnson, Mrs	female	27	0	2	347742	11.1333		S
11	10	1	2	Nasser, Mr	female	14	1	0	237736	30.0708		C
12	11	1	3	Sandstrom, Mrs	female	4	1	1	PP 9549	16.7	G6	S
13	12	1	1	Bonnell, Mrs	female	58	0	0	113783	26.55	C103	S
14	13	0	3	Saunders, Mr	male	20	0	0	A/5. 2151	8.05		S
15	14	0	3	Andersson, Mr	male	39	1	5	347082	31.275		S

target variable

# Classification and Deep Learning




# Euler's Function (e: 2.71828. . .)







# The sigmoid Activation Function



# The tanh Activation Function



# The ReLU Activation Function




# The softmax Activation Function

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}} \quad \text{for } j = 1, \dots, K.$$

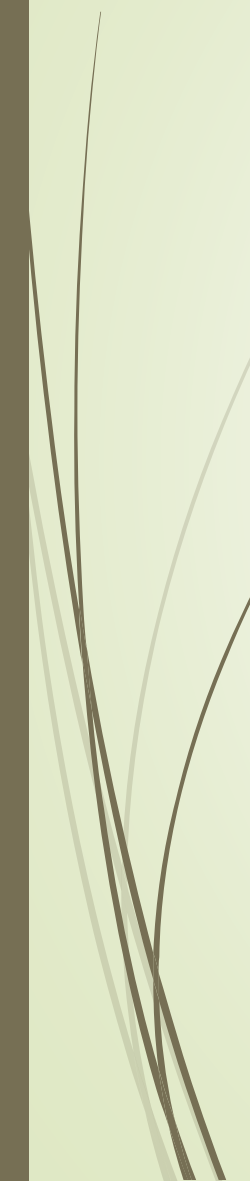


# Activation Functions in Python



```
import numpy as np
...
# Python sigmoid example:
z = 1 / (1 + np.exp(-np.dot(W, x)))
...
# Python tanh example:
z = np.tanh(np.dot(W, x));

# Python ReLU example:
z = np.maximum(0, np.dot(W, x))
```







# What's the “Best” Activation Function?

- Initially: **sigmoid** was popular
- Then: **tanh** became popular
- **Now**: **RELU** is preferred (better results)
- Softmax: for FC (fully connected) layers
- NB: **sigmoid** and **tanh** are used in LSTMs

# Linear Regression

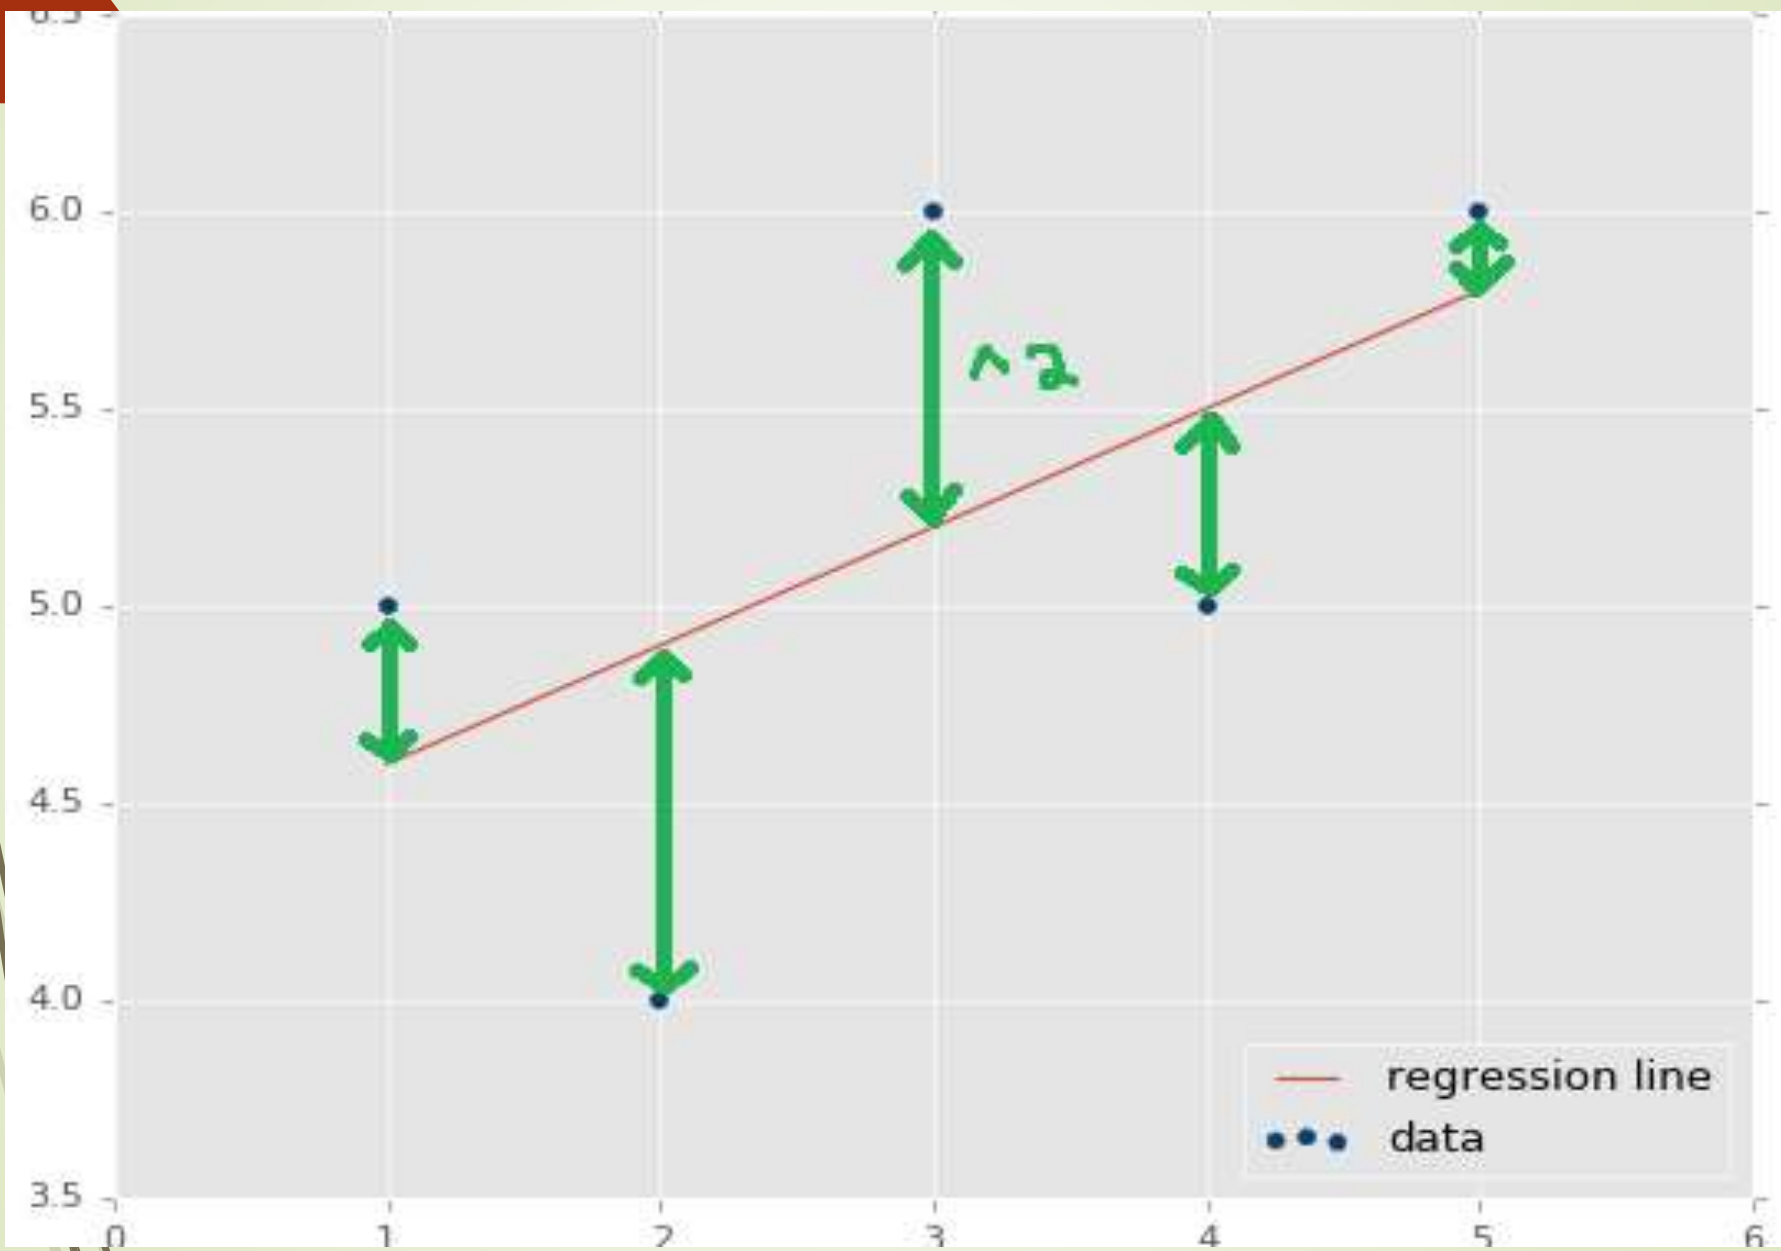
- ➡ One of the simplest models in ML
- ➡ Fits a line ( $y = m \cdot x + b$ ) to data in 2D
- ➡ Finds best line by minimizing MSE:  
 $m$  = slope of the best-fitting line  
 $b$  = y-intercept of the best-fitting line

# Which is Good for Linear Regression?







# Linear Regression in 2D: example




# Example #1: $MSE = ?$



## Example #2: MSE = ?




# Example #3 (Random Points)





# Sample Cost Function #1 (MSE)



# Linear Regression: example #1

- One feature (independent variable):

$X$  = number of square feet

- Predicted value (dependent variable):

$Y$  = cost of a house

- A very “coarse grained” model

- We can devise a much better model

# Linear Regression: example #2

- Multiple features:
- $X_1$  = # of square feet
- $X_2$  = # of bedrooms
- $X_3$  = # of bathrooms (dependency?)
- $X_4$  = age of house
- $X_5$  = cost of nearby houses
- $X_6$  = corner lot (or not): Boolean
- a much better model (6 features)

# Linear Multivariate Analysis

➡ General form of multivariate equation:

➡  $Y = w_1 * x_1 + w_2 * x_2 + \dots + w_n * x_n + b$

➡  $w_1, w_2, \dots, w_n$  are numeric values

➡  $x_1, x_2, \dots, x_n$  are variables (features)


➡ Properties of variables:

Can be independent (Naïve Bayes)


weak/strong dependencies can exist



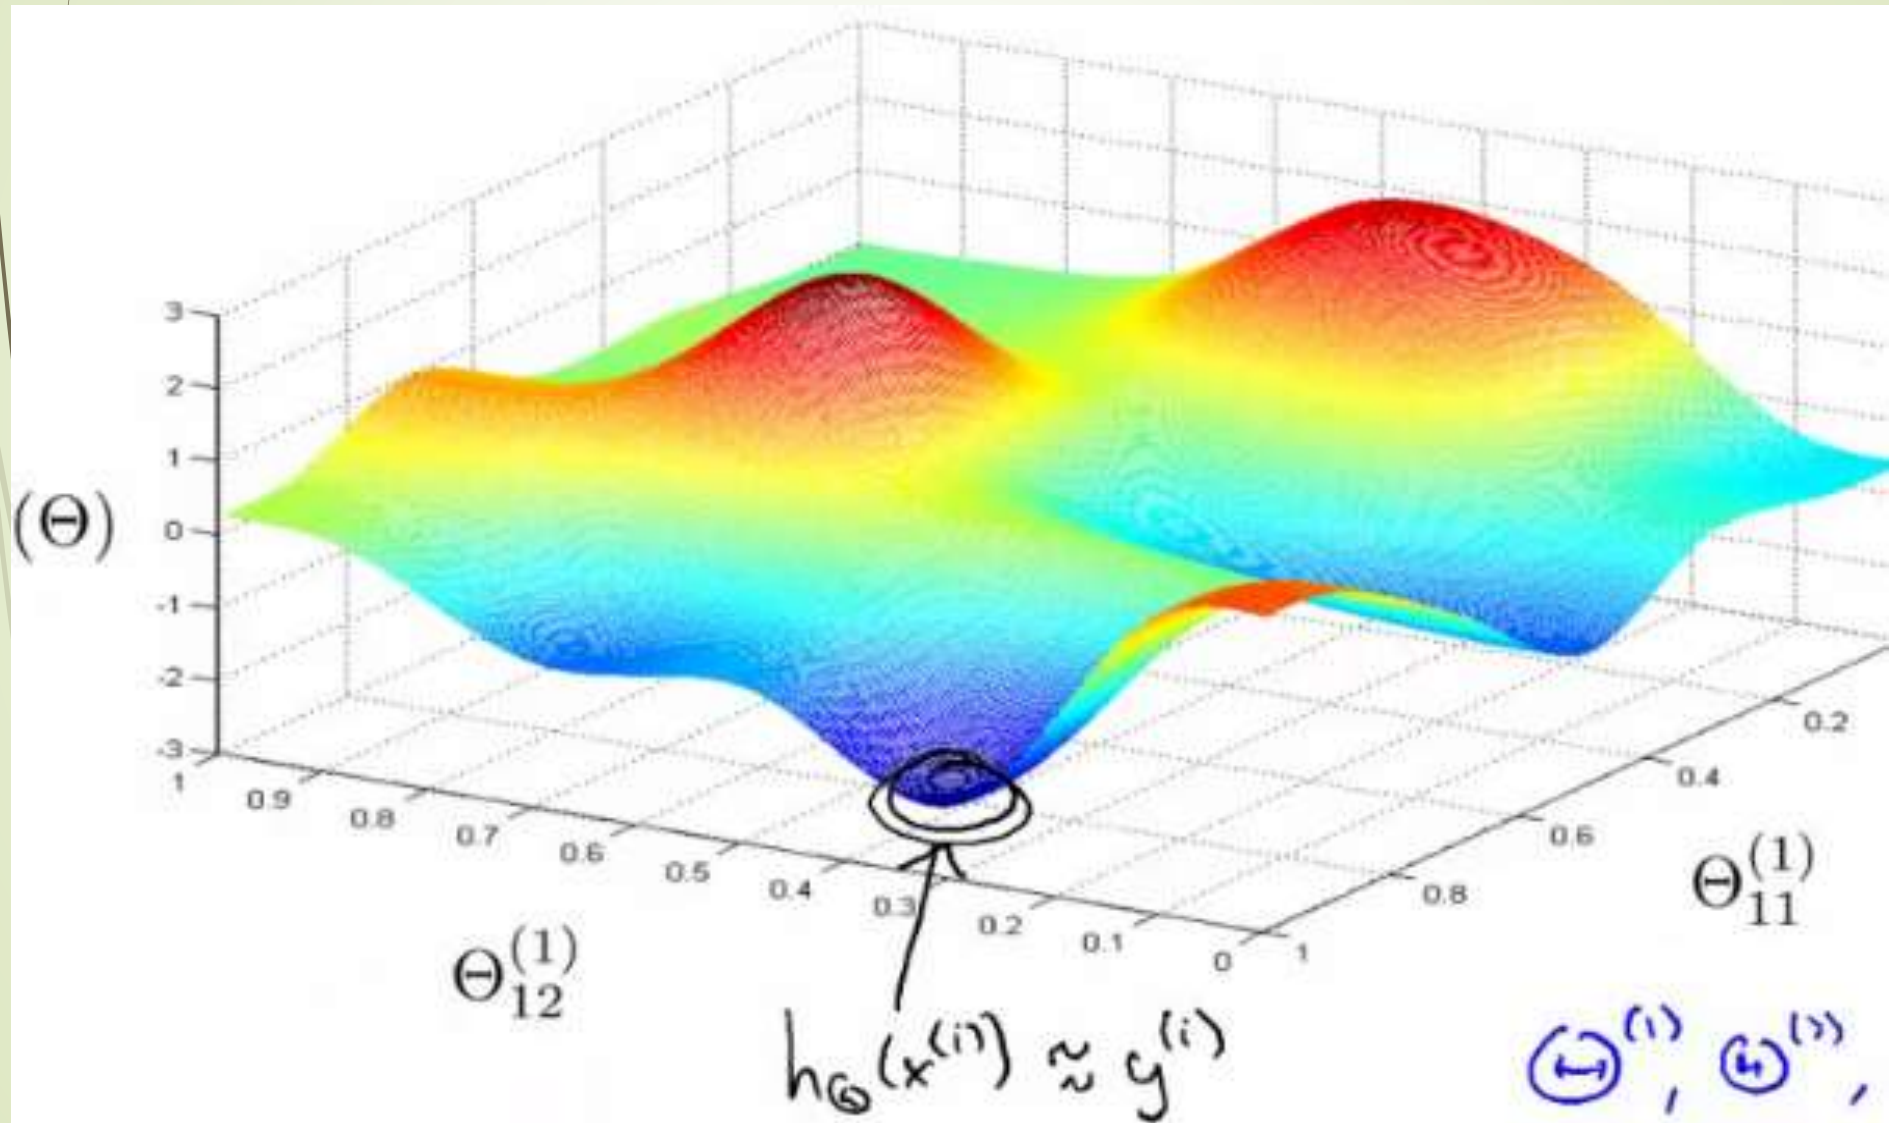
# Sample Cost Function #1 (MSE)



# Sample Cost Function #2



# Sample Cost Function #3





# Types of Optimizers

- SGD
- rmsprop
- Adagrad
- Adam
- Others

<http://cs229.stanford.edu/notes/cs229-notes1.pdf>

# Deep Neural Network: summary

- input layer, multiple hidden layers, and output layer
- nonlinear processing via activation functions
- perform transformation and feature extraction
- gradient descent algorithm with back propagation
- each layer receives the output from previous layer
- results are comparable/superior to human experts

# CNNs versus RNNs

## ➡ CNNs (Convolutional NNs):

Good for image processing

2000: CNNs processed 10-20% of all checks

=> Approximately 60% of all NNs

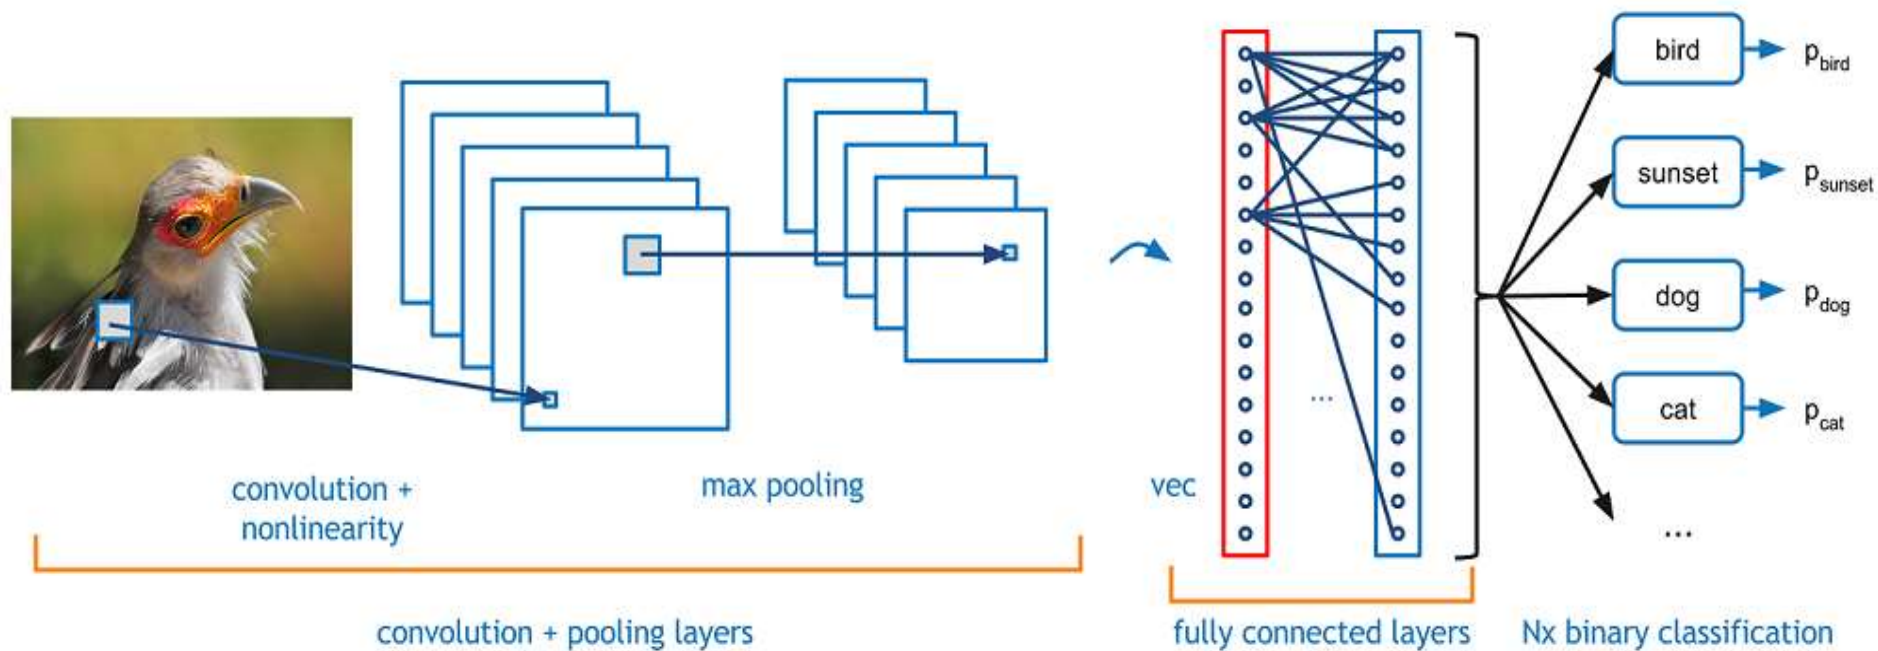
## ➡ RNNs (Recurrent NNs):

Good for NLP and audio

Used in hybrid networks



# CNNs: Convolution, ReLU, and Max Pooling



# CNNs: Convolution Calculations

35	40	41	45	50										
40	40	42	46	52										
42	46	50	55	55		0	1	0						
48	52	56	58	60		0	0	0						
56	60	65	70	75		0	0	0						

➡ <https://docs.gimp.org/en/plugin-convmatrix.html>

# CNNs: Convolution Matrices (examples)

➡ Sharpen:

0	0	0	0	0
0	0	-1	0	0
0	-1	5	-1	0
0	0	-1	0	0
0	0	0	0	0

➡ Blur:

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0

# CNNs: Convolution Matrices (examples)

➔ Edge detect:

	<b>0</b>	<b>1</b>	<b>0</b>	
	<b>1</b>	<b>-4</b>	<b>1</b>	
	<b>0</b>	<b>1</b>	<b>0</b>	

➔ Emboss:

	<b>-2</b>	<b>-1</b>	<b>0</b>	
	<b>-1</b>	<b>1</b>	<b>1</b>	
	<b>0</b>	<b>1</b>	<b>2</b>	

# CNNs: Max Pooling Example

Single depth slice

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters  
and stride 2

6	8
3	4

# GANs: Generative Adversarial Networks



"panda"

57.7% confidence

+  $\epsilon$



=



"gibbon"

99.3% confidence



# GANs: Generative Adversarial Networks

- Make imperceptible changes to images
- Can **consistently** defeat all NNs
- Can have **extremely** high error rate
- Some images create optical illusions
- <https://www.quora.com/What-are-the-pros-and-cons-of-using-generative-adversarial-networks-a-type-of-neural-network>

# GANs: Generative Adversarial Networks

## ► Create your own GANs:

<https://www.oreilly.com/learning/generative-adversarial-networks-for-beginners>

<https://github.com/jonbruner/generative-adversarial-networks>

## ► GANs from MNIST:

<http://edwardlib.org/tutorials/gan>

## ► GANs and Capsule networks?

# CNN in Python/Keras (fragment)

```
from keras.models import Sequential
```

```
from keras.layers.core import Dense, Dropout, Activation
```

```
from keras.layers.convolutional import Conv2D, MaxPooling2D
```

```
from keras.optimizers import Adadelta
```

```
input_shape = (3, 32, 32)
```

```
nb_classes = 10
```

```
model = Sequential()
```

```
model.add(Conv2D(32,(3, 3),padding='same',  
                 input_shape=input_shape))
```

```
model.add(Activation('relu'))
```

```
model.add(Conv2D(32, (3, 3)))
```

```
model.add(Activation('relu'))
```

```
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
model.add(Dropout(0.25))
```

# What is TensorFlow?

- An open source framework for ML and DL
- A “computation” graph
- Created by Google (released 11/2015)
- Evolved from Google Brain
- Linux and Mac OS X support (VM for Windows)
- TF home page: <https://www.tensorflow.org/>

# What is TensorFlow?

- Support for Python, Java, C++
- Desktop, server, mobile device (TensorFlow Lite)
- CPU/GPU/TPU support
- Visualization via TensorBoard
- Can be embedded in Python scripts
- Installation: `pip install tensorflow`

TensorFlow cluster:

<https://www.tensorflow.org/deploy/distributed>

# TensorFlow Use Cases (Generic)

- Image recognition
- Computer vision
- Voice/sound recognition
- Time series analysis
- Language detection
- Language translation
- Text-based processing
- Handwriting Recognition



# Aspects of TensorFlow

- Graph: graph of operations (DAG)
- Sessions: contains Graph(s)
- lazy execution (default)
- operations in parallel (default)
- Nodes: operators/variables/constants
- Edges: tensors
- => graphs are split into subgraphs and executed in parallel (or multiple CPUs)

# TensorFlow Graph Execution

- ➡ Execute statements in a `tf.Session()` object
- ➡ Invoke the “run” method of that object
- ➡ “eager” execution is available ( $\geq$  v1.4)
- ➡ included in the mainline (v1.7)
- ➡ Installation: `pip install tensorflow`

# What is a Tensor?

- TF tensors are n-dimensional arrays
- TF tensors are very similar to numpy ndarrays
- **scalar number:** a zeroth-order tensor
- **vector:** a first-order tensor
- **matrix:** a second-order tensor
- **3-dimensional array:** a 3rd order tensor
- <https://dzone.com/articles/tensorflow-simplified-examples>

# TensorFlow “primitive types”

## ➡ `tf.constant`:

- + initialized immediately
- + immutable

## ➡ `tf.placeholder` (a function):

- + initial value is not required
- + can have variable shape
- + assigned value via `feed_dict` at run time
- + receive data from “external” sources

# TensorFlow “primitive types”

- ➡ **tf.Variable** (a class):
  - + initial value is required
  - + updated during training
  - + maintain state across calls to “run()”
  - + in-memory buffer (saved/restored from disk)
  - + can be shared in a distributed environment
  - + they hold learned parameters of a model

# TensorFlow: constants (immutable)

➤ import tensorflow as tf

➤ aconst = tf.constant(3.0)

➤ print(aconst)

# output: Tensor("Const:0", shape=(), dtype=float32)

➤ sess = tf.Session()

➤ print(sess.run(aconst))

# output: 3.0

➤ sess.close()

➤ # => there's a better way 😊



# TensorFlow: constants

- `import tensorflow as tf`
- `aconst = tf.constant(3.0)`
- `print(aconst)`
- **Automatically close “sess”**
- **with** `tf.Session() as sess:`
- `print(sess.run(aconst))`

# TensorFlow Arithmetic

```
import tensorflow as tf
```

```
a = tf.add(4, 2)
```

```
b = tf.subtract(8, 6)
```

```
c = tf.multiply(a, 3)
```

```
d = tf.div(a, 6)
```

```
with tf.Session() as sess:
```

```
    print(sess.run(a)) # 6
```

```
    print(sess.run(b)) # 2
```

```
    print(sess.run(c)) # 18
```

```
    print(sess.run(d)) # 1
```

# TF placeholders and feed\_dict


```
import tensorflow as tf

a = tf.placeholder("float")
b = tf.placeholder("float")
c = tf.multiply(a,b)

# initialize a and b:
feed_dict = {a:2, b:3}

# multiply a and b:
with tf.Session() as sess:
    print(sess.run(c, feed_dict))
```

# TensorFlow: Simple Equation



```
import tensorflow as tf

# W and x are 1d arrays
W = tf.constant([10,20], name='W')
X = tf.placeholder(tf.int32, name='x')
b = tf.placeholder(tf.int32, name='b')

Wx = tf.multiply(W, x, name='Wx')
y = tf.add(Wx, b, name='y') OR
y2 = tf.add(tf.multiply(W,x),b)
```

# TensorFlow fetch/feed\_dict

```
with tf.Session() as sess:
    print("Result 1: Wx = ",
          sess.run(Wx, feed_dict={x: [5, 10]}))

    print("Result 2: y = ",
          sess.run(y, feed_dict={x: [5, 10], b: [15, 25]}))
```

➡ Result 1: Wx = [50 200]

➡ Result 2: y = [65 225]

# Saving Graphs for TensorBoard

```
import tensorflow as tf

x = tf.constant(5, name="x")
y = tf.constant(8, name="y")
z = tf.Variable(2*x+3*y, name="z")
init = tf.global_variables_initializer()

with tf.Session() as session:
    writer = tf.summary.FileWriter("./tf_logs", session.graph)
    session.run(init)
    print 'z = ', session.run(z) # => z = 34

# launch: tensorboard -logdir=./tf_logs
```



# TensorFlow Eager Execution

- An imperative interface to TF
- Fast debugging & immediate run-time errors
- Eager execution is “mainline” in v1.7 of TF
- => requires Python 3.x (not Python 2.x)

# TensorFlow Eager Execution

- integration with Python tools
- Supports dynamic models + Python control flow
- support for custom and higher-order gradients
- Supports **most** TensorFlow operations
- => Default mode in TensorFlow 2.0 (2019)
- <https://research.googleblog.com/2017/10/eager-execution-imperative-define-by.html>

# TensorFlow Eager Execution

- `import tensorflow as tf`
- `import tensorflow.contrib.eager as tfe`
- `tfe.enable_eager_execution()`
- `x = [[2.]]`
- `m = tf.matmul(x, x)`
- `print(m)`
- `# tf.Tensor([[4.]], shape=(1, 1), dtype=float32)`

# What is tensorflow.js?

- an ecosystem of JS tools for machine learning
- TensorFlow.js also includes a Layers API
- a library for building machine learning models
- tools to port TF SavedModels & Keras HDF5 models
- => <https://js.tensorflow.org/>

# What is tensorflow.js?

- tensorflow.js evolved from deeplearn.js
- deeplearn.js is now called TensorFlow.js Core
- TensorFlow.js Core: a flexible low-level API
- TensorFlow.js Layers:  
a high-level API similar to Keras
- TensorFlow.js Converter:  
tools to import a TF SavedModel to TensorFlow.js

# async keyword

- keyword placed before JS functions
- For functions that return a Promise
- Trivial example:

```
async function f() {  
    return 1;  
}
```



# await keyword

- Works only inside async JS functions
- Trivial example:
- `let value = await mypromise;`

# async/await example

```
async function f() {  
  let promise = new Promise((resolve, reject) => {  
    setTimeout(() => resolve("done!"), 1000)  
  });  
  
  // wait till the promise resolves  
  let result = await promise  
  alert(result)  
}
```

f()

# Tensorflow.js Samples

- 1) `tfjs-example.html` (linear regression)
- 2) `js.tensorflow.org` (home page)
- 3) `https://github.com/tensorflow/tfjs-examples-master`
  - a) `cd mnist-core`
  - b) `yarn`
  - c) `yarn watch`

# Deep Learning and Art/"Stuff"

➡ "Convolutional Blending" images:

=> 19-layer Convolutional Neural Network

[www.deepart.io](http://www.deepart.io)

➡ <https://www.fastcodesign.com/90124942/this-google-engineer-taught-an-algorithm-to-make-train-footage-and-its-hypnotic>