


Introduction to TensorFlow 2



Oswald Campesato
ocampesato@yahoo.com

Highlights/Overview

- What is TensorFlow 2?
- Major Changes in TF 2
- Working with strings/arrays/tensors
- Working with `@tf.function` decorator
- Working with generators
- Working with `tf.data.Dataset`
- Datasets in TF 1.x versus TF 2
- Working with `tf.keras`
- CNNs, RNNs, LSTMs, Bidirectional LSTMs
- Reinforcement Learning

What is TensorFlow 2?

- An open source framework for ML and DL
- Created by Google (TF 1.x released 11/2015)
- Evolved from Google Brain
- Linux and Mac OS X support (VM for Windows)
- TF home page: <https://www.tensorflow.org/>

Status of TensorFlow 2



- Alpha release: around March/2019

- Other “Nightly” builds:

<https://pypi.org/project/tf-nightly-2.0-preview/2.0.0.dev20190417/#files>

- Production release: TBD (later this year)

TF 2 Platform Support

- Tested and supported on 64-bit systems
- Ubuntu 16.04 or later
- Windows 7 or later
- macOS 10.12.6 (Sierra) or later (no GPU support)
- Raspbian 9.0 or later

TF 2 Docker Images



- <https://hub.docker.com/r/tensorflow/tensorflow/>

```
$ docker run -it --rm tensorflow/tensorflow bash
```

- Start a CPU-only container with Python 2:

```
$ docker run -it --rm --runtime=nvidia tensorflow/tensorflow:latest-gpu-py3 python
```

- Start a GPU container with Python 3 & Python interpreter:

```
$ docker run -it --rm -v $(realpath ~/notebooks):/tf/notebooks -p 8888:8888 tensorflow/tensorflow:late
```

Status of TensorFlow 1.x



- 1.13.1: current stable release

- 1.14: latest release

- After TF 2 is in production:

Only security-related updates to TF 1.x

TF 1.x support for one year after TF 2 production

What is TensorFlow 2?

- Support for Python, Java, C++
- Desktop, Server, Mobile, Web
- CPU/GPU/TPU support
- Visualization via TensorBoard
- Can be embedded in Python scripts
- Installation: `pip install tensorflow`
- Ex: `pip install tensorflow==2.0.0-alpha0`

TensorFlow Use Cases (Generic)

- 
- Image recognition
 - Computer vision
 - Voice/sound recognition
 - Time series analysis
 - Language detection
 - Language translation
 - Text-based processing
 - Handwriting Recognition
- 

Major Changes in TF 2

- TF 2 Eager Execution: default mode
- `@tf.function` decorator: instead of `tf.Session()`
- AutoGraph: graph generated in `@tf.function()`
- Generators: used with `tf.data.Dataset` (TF 2)
- Differentiation: calculates gradients
- `tf.GradientTape`: automatic differentiation

Removed from TensorFlow 2

- `tf.Session()`
- `tf.placeholder()`
- `tf.global_initializer()`
- `feed_dict`
- Variable scopes
- `tf.contrib` code
- `tf.flags` and `tf.contrib`
- Global variables

- => TF 1.x functions moved to `compat.v1`

Upgrading Files/Directories to TF 2

- The TF 2 upgrade script:

<https://www.tensorflow.org/alpha/guide/upgrade>

- Script name: `tf_upgrade_v2`

- install TF to get the upgrade script:

`pip install tensorflow (NOT pip3)`

Upgrading Files/Directories to TF 2

- 1) upgrade one file:

```
tf_upgrade_v2 --infile oldtf.py --outfile newtf.py
```

- The preceding command creates **report.txt**

- Upgrade an entire directory:


```
tf_upgrade_v2 --intree mydir --outtree newtf.py --  
copyotherfiles False
```

- **NB: do NOT make manual upgrade changes**

Migrating to TensorFlow 2

- 1) do not manually update parts of your code
 - 2) script won't work with functions with reordered arguments
 - 3) script assumes this TF statement: `import tensorflow as tf`
 - 4) the script does not reorder arguments
 - 5) the script adds keyword arguments to functions that have their arguments reordered.
- Upgrade Jupyter notebooks and Python files in a GitHub repo:
<http://tf2up.ml/>
- ⇒ Replace prefix <https://github.com/> with <http://tf2up.m>

Check your TF Version



```
➤ import tensorflow as tf
➤ import numpy

➤ print("TF version: ",tf.__version__)
➤ print("Keras version:",tf.keras.__version__)
➤ print("Numpy version:",numpy.__version__)

➤ # my system:
➤ #TF version:      2.0.0-alpha0
➤ #Keras version:  2.2.4-tf
➤ #Numpy version:  1.15.4
```

What is a Tensor?

- ➡ TF tensors are n-dimensional arrays
- ➡ TF tensors are very similar to numpy ndarrays
- ➡ **scalar number:** a zeroth-order tensor
- ➡ **vector:** a first-order tensor
- ➡ **matrix:** a second-order tensor
- ➡ **3-dimensional array:** a 3rd order tensor
- ➡ <https://dzone.com/articles/tensorflow-simplified-examples>

TensorFlow Eager Execution

- An imperative interface to TF
- Fast debugging & immediate run-time errors
- => TF 2 requires Python 3.x (not Python 2.x)
- No static graphs or sessions

TensorFlow Eager Execution

- integration with Python tools
- Supports dynamic models + Python control flow
- support for custom and higher-order gradients
- => Default mode in TensorFlow 2.0

TensorFlow Eager Execution

```
import tensorflow as tf

x = [[2.]]

# m = tf.matmul(x, x)  <= deprecated
m = tf.linalg.matmul(x, x)

print("m:", m)
print("m:", m.numpy())
```

#Output:

```
#m: tf.Tensor([[4.]], shape=(1,1), dtype=float32)
#m: [[4.]]
```

Basic TF 2 Operations

- Working with Constants
- Working with Strings
- Working with Tensors
- Working Operators
- Arrays in TF 2
- Multiplying Two Arrays
- Convert Python Arrays to TF Arrays
- Conflicting Types

TensorFlow “primitive types”

➤ `tf.constant`:

> initialized immediately and immutable

➤ `import tensorflow as tf`

➤ `aconst = tf.constant(3.0)`

➤ `print(aconst)`

output for TF 2:

`#tf.Tensor(3.0, shape=(), dtype=float32)`

output for TF 1.x:

`#Tensor("Const:0", shape=(), dtype=float32)`

TensorFlow “primitive types”


- ➡ **tf.Variable** (a class):
 - + initial value is required
 - + updated during training
 - + in-memory buffer (saved/restored from disk)
 - + can be shared in a distributed environment
 - + they hold learned parameters of a model

TensorFlow Arithmetic

```
import tensorflow as tf # arith1.py  
a = tf.add(4, 2)  
b = tf.subtract(8, 6)  
c = tf.multiply(a, 3)  
d = tf.div(a, 6)
```

```
➤ print(a) # 6  
➤ print(b) # 2  
➤ print(c) # 18  
➤ print(d) # 1  
➤ print("a:", a.numpy())  
➤ print("b:", b.numpy())  
➤ print("c:", c.numpy())  
➤ print("d:", d.numpy())
```

TensorFlow Arithmetic



- `tf.Tensor(6, shape=(), dtype=int32)`
- `tf.Tensor(2, shape=(), dtype=int32)`
- `tf.Tensor(18, shape=(), dtype=int32)`
- `tf.Tensor(1.0, shape=(), dtype=float64)`



- 6

- 2

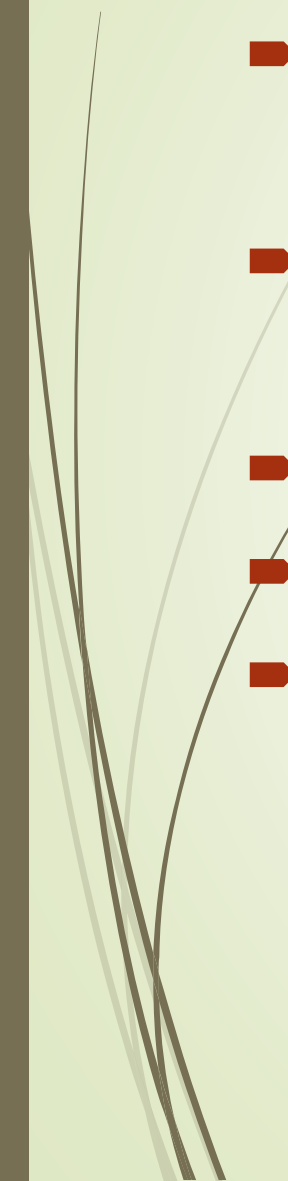

- 18

- 1.0

TF 2 Loops and Arrays



- Using “for” loops
- Using “while” loops
- `tf.reduce_prod()`
- `tf.reduce_sum()`

A “for” Loop Example




- `import tensorflow as tf`
- `x = tf.Variable(0, name='x')`
- `for i in range(5):`
 - `x = x + 1`
 - `print("x:", x)`

A “for” Loop Example



- `x: tf.Tensor(1, shape=(), dtype=int32)`
- `x: tf.Tensor(2, shape=(), dtype=int32)`
- `x: tf.Tensor(3, shape=(), dtype=int32)`
- `x: tf.Tensor(4, shape=(), dtype=int32)`
- `x: tf.Tensor(5, shape=(), dtype=int32)`

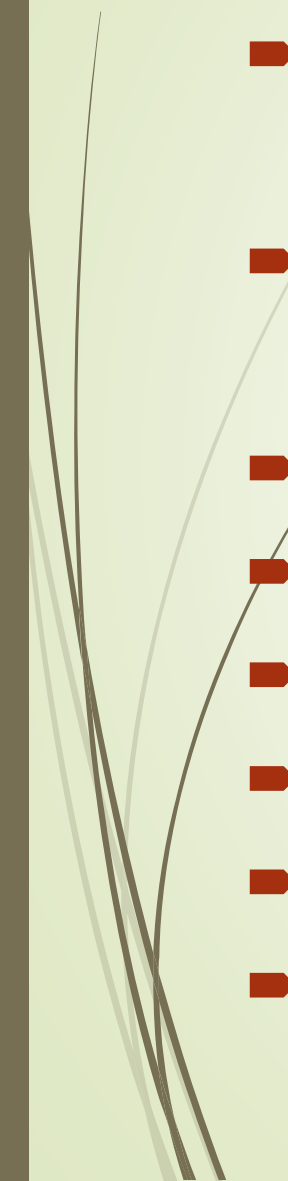
A “while” Loop Example



```
➤ import tensorflow as tf

➤ a = tf.constant(12)

➤ while not tf.equal(a, 1):
➤     if tf.equal(a % 2, 0):
➤         a = a / 2
➤     else:
➤         a = 3 * a + 1
➤     print(a)
```



A “while” Loop Example

- `tf.Tensor(6.0, shape=(), dtype=float64)`
- `tf.Tensor(3.0, shape=(), dtype=float64)`
- `tf.Tensor(10.0, shape=(), dtype=float64)`
- `tf.Tensor(5.0, shape=(), dtype=float64)`
- `tf.Tensor(16.0, shape=(), dtype=float64)`
- `tf.Tensor(8.0, shape=(), dtype=float64)`
- `tf.Tensor(4.0, shape=(), dtype=float64)`
- `tf.Tensor(2.0, shape=(), dtype=float64)`
- `tf.Tensor(1.0, shape=(), dtype=float64)`

Working with Arrays

- `import tensorflow as tf`
- `import numpy as np`
- `# create a Python array:`
- `array_1d = np.array([1.3, 1, 4.0, 23.5])`
- `tf_tensor =`
`tf.convert_to_tensor(value=array_1d,`
`dtype=tf.float64)`
- `print(tf_tensor)`
- `print("0:",tf_tensor[0])`
- `print("2:",tf_tensor[2])`

Working with Arrays

- `tf.Tensor([1.3 1. 4. 23.5],
shape=(4,), dtype=float64)`
- `0: tf.Tensor(1.3, shape=(), dtype=float64)`
- `2: tf.Tensor(4.0, shape=(), dtype=float64)`

Random Numbers

- The `tf.random.normal()`
- The `tf.truncated_normal()`
- TF Arrays with Random Values

Random Numbers: examples


```
import tensorflow as tf

# normal distribution:
w=tf.Variable(tf.random.normal([784,10],stddev=0.01))

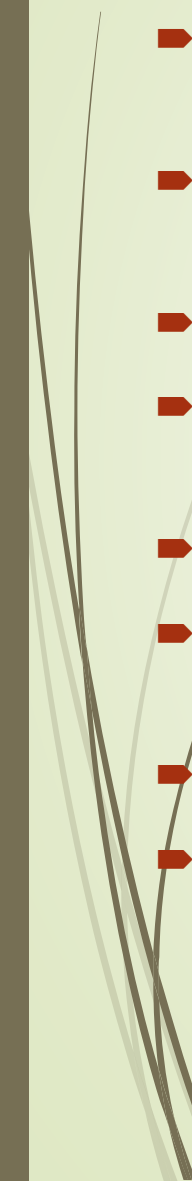
# mean of an array:
b = tf.Variable([10,20,30,40,50,60],name='t')

print("w: ",w)
print("b: ",tf.reduce_mean(input_tensor=b))
```

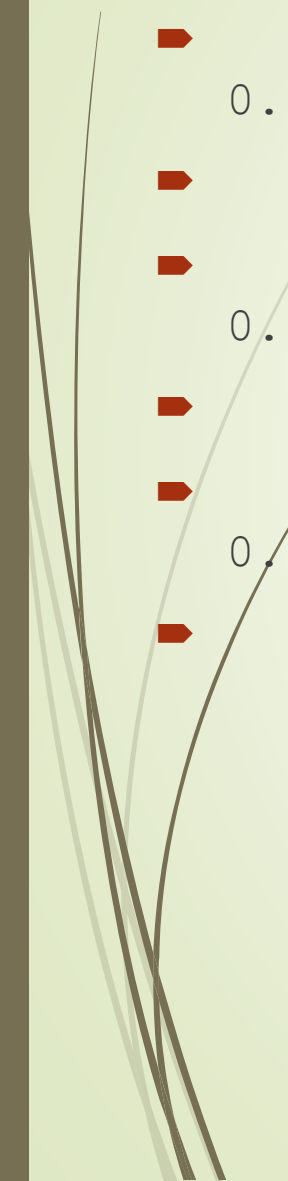

Random Numbers: examples



```
w: <tf.Variable 'Variable:0' shape=(784, 10)
dtype=float32, numpy=
array([[ 0.00199239,  0.00285635,  0.00804297, ..., -
 0.00323935,
        0.00138759,  0.00941323],
       [-0.00017284, -0.00708508, -0.00670188, ...,
 0.00085814,
        -0.01298123,  0.01133613],
       [-0.00506489,  0.01542902, -0.00710952, ..., -
 0.00294803,
        -0.00767813,  0.00815126],
       ...,
```


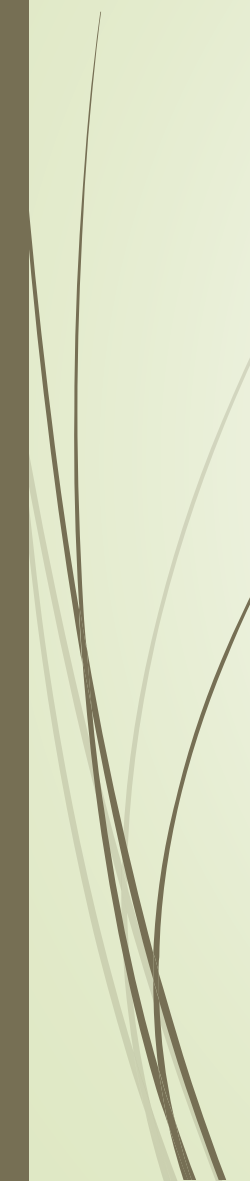


Random Numbers: examples

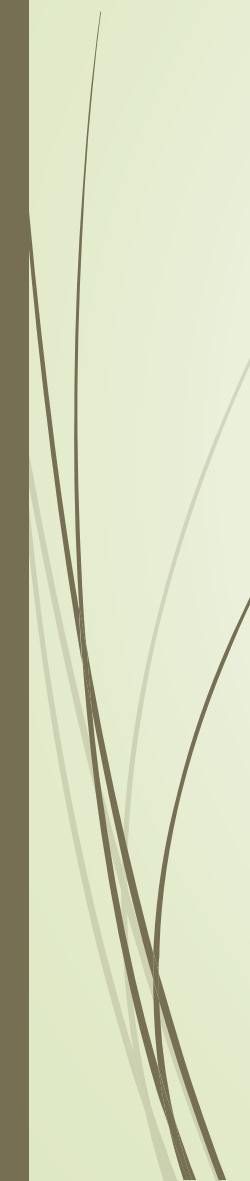



```
[ 0.01556268,  0.01296226,  0.01230366, ...,  
0.01154588,  
-0.01639041,  0.00107052],  
[ 0.0109421 ,  0.00486962,  0.02887715, ..., -  
0.00792963,  
-0.00337163,  0.0041892 ],  
[-0.0199652 , -0.00471972,  0.00246108, ..., -  
0.00995417,  
-0.00854415, -0.01487656]], dtype=float32)>
```

Useful TF 2 APIs

- 
- `tf.shape()`
 - `tf.rank()`
 - `tf.range()`
 - `tf.reshape()`
 - `tf.ones()`
 - `tf.zeros()`
 - `tf.fill()`
- 

The tf.range() API



- `import tensorflow as tf`
- `a1 = tf.range(3, 18, 3)`
- `a2 = tf.range(0, 8, 2)`
- `a3 = tf.range(-6, 6, 3)`
- `a4 = tf.range(-10, 10, 4)`
- `print('a1:', a1)`
- `print('a2:', a2)`
- `print('a3:', a3)`
- `print('a4:', a4)`

The tf.range() API

```
a1: tf.Tensor([3 6 9 12 15], shape=(5,), dtype=int32)

a2: tf.Tensor([0 2 4 6], shape=(4,), dtype=int32)

a3: tf.Tensor([-6 -3 0 3], shape=(4,), dtype=int32)

a4: tf.Tensor([-10 -6 -2 2 6], shape=(5,), dtype=int32)
```

Some TF 2 “lazy operators”

- `map()`
- `filter()`
- `flatmap()`
- `batch()`
- `take()`
- `zip()`
- `flatten()`

- Combined via “method chaining”

TF 2 “lazy operators”

- `filter()`:
 - uses Boolean logic to "filter" the elements in an array to determine which elements satisfy the Boolean condition
- `map()`: a projection
 - this operator "applies" a lambda expression to each input element
- `flat_map()`:
 - maps a single element of the input dataset to a Dataset of elements

TF 2 “lazy operators”

- `batch(n)`:
 - processes a "batch" of n elements during each iteration
- `repeat(n)`:
 - repeats its input values n times
- `take(n)`:
 - operator "takes" n input values

TF 2 tf.data.Dataset

- TF “wrapper” class around data sources
- Located in the tf.data.Dataset namespace
- Supports lambda expressions
- Supports lazy operators
- They use **generators** instead of **iterators** (TF 1.x)

tf.data.Dataset.from_tensors()

- Import tensorflow as tf
- #combine the input into one element
- `t1 = tf.constant([[1, 2], [3, 4]])`
- `ds1 = tf.data.Dataset.from_tensors(t1)`
- # output: `[[1, 2], [3, 4]]`

tf.data.Dataset.from_tensor_slices()

- Import tensorflow as tf
- #separate element for each item
- t2 = tf.constant([[1, 2], [3, 4]])
- ds1 = tf.data.Dataset.from_tensor_slices(t2)
- # output: [1, 2], [3, 4]

TF 2 Datasets: code sample

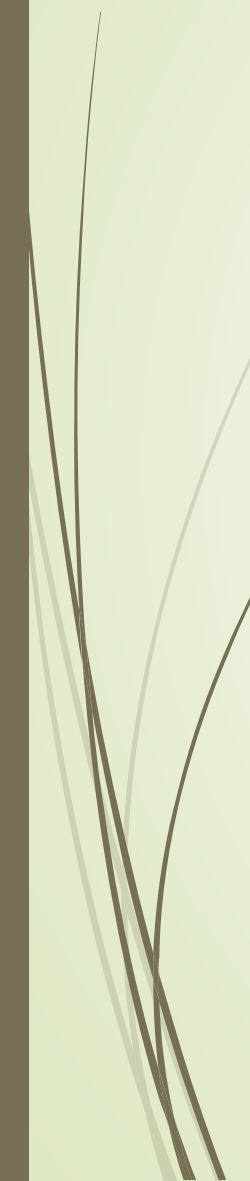

- `import tensorflow as tf # tf2-dataset.py`
- `import numpy as np`

```
x = np.arange(0, 10)
```

- `# create a dataset from a Numpy array`

```
ds = tf.data.Dataset.from_tensor_slices(x)
```

TF 1.x Datasets: iterator



```
➤ import tensorflow as tf # tf1x-plusone.py
➤ import numpy as np

➤ x = np.arange(0, 10)
➤ # create dataset object from Numpy array
➤ ds = tf.data.Dataset.from_tensor_slices(x)
➤ ds.map(lambda x: x + 1)

➤ # create a one-shot iterator <= TF 1.x
➤ iterator = ds.make_one_shot_iterator()

➤ for i in range(10):
➤     val = iterator.get_next()
➤     print("val:", val)
```

TF 2 Datasets: generator

```
import tensorflow as tf # tf2-plusone.py
import numpy as np
x = np.arange(0, 5) # 0, 1, 2, 3, 4

def gener():
    for i in x:
        yield (3*i)

ds = tf.data.Dataset.from_generator(gener, (tf.int64))


for value in ds.take(len(x)):
    print("1value:",value)

for value in ds.take(2*len(x)):
    print("2value:",value)
```

TF 2 Datasets: generator

- 1value: tf.Tensor(0, shape=(), dtype=int64)
- 1value: tf.Tensor(3, shape=(), dtype=int64)
- 1value: tf.Tensor(6, shape=(), dtype=int64)
- 1value: tf.Tensor(9, shape=(), dtype=int64)
- 1value: tf.Tensor(12, shape=(), dtype=int64)
- 2value: tf.Tensor(0, shape=(), dtype=int64)
- 2value: tf.Tensor(3, shape=(), dtype=int64)
- 2value: tf.Tensor(6, shape=(), dtype=int64)
- 2value: tf.Tensor(9, shape=(), dtype=int64)
- 2value: tf.Tensor(12, shape=(), dtype=int64)

TF 1.x map() and take()



```
➤ import tensorflow as tf # tf 1.12.0
➤ tf.enable_eager_execution()

➤ ds = tf.data.TextLineDataset("file.txt")
➤ ds = ds.map(lambda line:
    tf.string_split([line]).values)

➤ try:
➤     for value in ds.take(2):
➤         print("value:", value)
➤ except tf.errors.OutOfRangeError:
➤     pass
```



TF 2 map() and take(): [file.txt](#)

- this is file line #1
- this is file line #2
- this is file line #3
- this is file line #4
- this is file line #5

TF 2 map() and take(): output

- ('value:', <tf.Tensor: id=16, shape=(5,), dtype=string, numpy=array(['this', 'is', 'file', 'line', '#1'], dtype=object)>)
- ('value:', <tf.Tensor: id=18, shape=(5,), dtype=string, numpy=array(['this', 'is', 'file', 'line', '#2'], dtype=object)>)

TF 2 map() and take()



```
import tensorflow as tf
import numpy as np

x = np.array([[1], [2], [3], [4]])

# make a ds from a numpy array
ds = tf.data.Dataset.from_tensor_slices(x)
ds = ds.map(lambda x: x*2).map(lambda x:
x+1).map(lambda x: x**3)

for value in ds.take(4):
    print("value:", value)
```

TF 2 map() and take(): output

- value: tf.Tensor([27], shape=(1,), dtype=int64)
- value: tf.Tensor([125], shape=(1,), dtype=int64)
- value: tf.Tensor([343], shape=(1,), dtype=int64)
- value: tf.Tensor([729], shape=(1,), dtype=int64)

TF 2 batch(): EXTRA CREDIT

```
import tensorflow as tf
```

```
import numpy as np
```

```
x = np.arange(0, 12)
```

```
def gener():
```

```
    i = 0
```

```
    while(i < len(x/3)):
```

```
        yield (i, i+1, i+2) # three integers at a time
```

```
        i += 3
```

```
ds = tf.data.Dataset.from_generator(gener, (tf.int64,tf.int64,tf.int64))
```

```
third = int(len(x)/3)
```

```
for value in ds.take(third):
```

```
    print("value:",value)
```

TF 2 batch() & zip(): EXTRA CREDIT

```
import tensorflow as tf
```

```
ds1 = tf.data.Dataset.range(100)
```

```
ds2 = tf.data.Dataset.range(0, -100, -1)
```


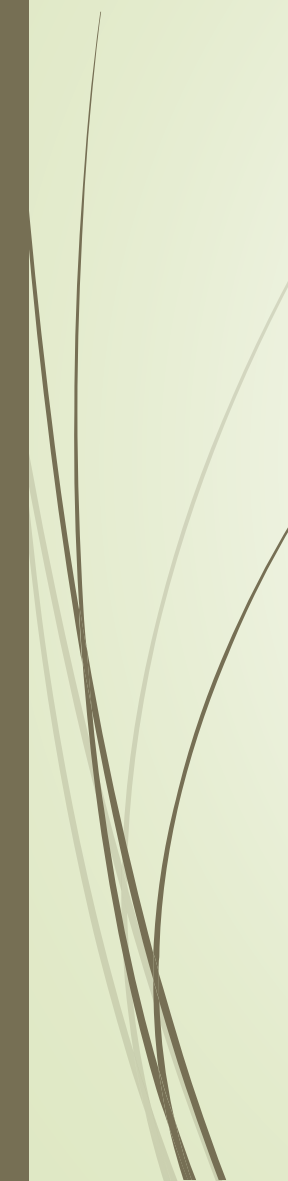
```
ds3 = tf.data.Dataset.zip((ds1, ds2))
```

```
ds4 = ds3.batch(4)
```

```
for value in ds4.take(10):
```

```
    print("value:", value)
```

Working with tf.keras

- 
- `tf.keras.layers`
 - `tf.keras.models`
 - `tf.keras.optimizers`
 - `tf.keras.utils`
 - `tf.keras.regularizers`
- 

Working with `tf.keras.models`

- `tf.keras.models.Sequential()`: most common
- `clone_model(...)`: Clone any Model instance
- `load_model(...)`: Loads a model saved via `save_model`
- `model_from_config(...)`: Instantiates a Keras model from its config
- `model_from_json(...)`:
 - Parses a JSON model config file and returns a model instance
- `model_from_yaml(...)`:
 - Parses a yaml model config file and returns a model instance
- `save_model(...)`: Saves a model to a HDF5 file

Working with tf.keras.layers

- tf.keras.models.Sequential()
- tf.keras.layers.Conv2D()
- tf.keras.layers.MaxPooling2D()
- tf.keras.layers.Flatten()
- tf.keras.layers.Dense()
- tf.keras.layers.Dropout(0.1)
- tf.keras.layers.BatchNormalization()
- tf.keras.layers.embedding()
- tf.keras.layers.RNN()
- tf.keras.layers.LSTM()
- tf.keras.layers.Bidirectional (ex: BERT)

TF 2 Activation Functions

- `tf.keras.activations.relu`
- `tf.keras.activations.selu`
- `tf.keras.activations.elu`
- `tf.keras.activations.linear`
- `tf.keras.activations.sigmoid`
- `tf.keras.activations.softmax`
- `tf.keras.activations.softplus`
- `tf.keras.activations.tanh`
- Others ...

TF 2 Datasets

- `tf.keras.datasets.boston_housing`
- `tf.keras.datasets.cifar10`
- `tf.keras.datasets.cifar100`
- `tf.keras.datasets.fashion_mnist`
- `tf.keras.datasets.imdb`
- `tf.keras.datasets.mnist`
- `tf.keras.datasets.reuters`

TF 2 “Experimental”

- `tf.keras.experimental.CosineDecay`
- `tf.keras.experimental.CosineDecayRestarts`
- `tf.keras.experimental.LinearCosineDecay`
- `tf.keras.experimental.NoisyLinearCosineDecay`
- `tf.keras.experimental.PeepholeLSTMCell`

TF 2 Optimizers

- `tf.optimizers.Adadelta`
- `tf.optimizers.Adagrad`
- `tf.optimizers.Adam`
- `tf.optimizers.Adamax`
- `tf.optimizers.SGD`

TF 2 Callbacks

- `tf.keras.callbacks.Callback`
- `tf.keras.callbacks.EarlyStopping`
- `tf.keras.callbacks.LambdaCallback`
- `tf.keras.callbacks.ModelCheckpoint`
- `tf.keras.callbacks.TensorBoard`
- `tf.keras.callbacks.TerminateOnNan`

TF 2 Losses

- `tf.losses.BinaryCrossEntropy`
- `tf.losses.CategoricalCrossEntropy`
- `tf.losses.CategoricalHinge`
- `tf.losses.CosineSimilarity`
- `tf.losses.Hinge`
- `tf.losses.Huber`
- `tf.losses.KLD`
- `tf.losses.KLDivergence`
- `tf.losses.MAE`
- `tf.losses.MSE`
- `tf.losses.SparseCategoricalCrossentropy`

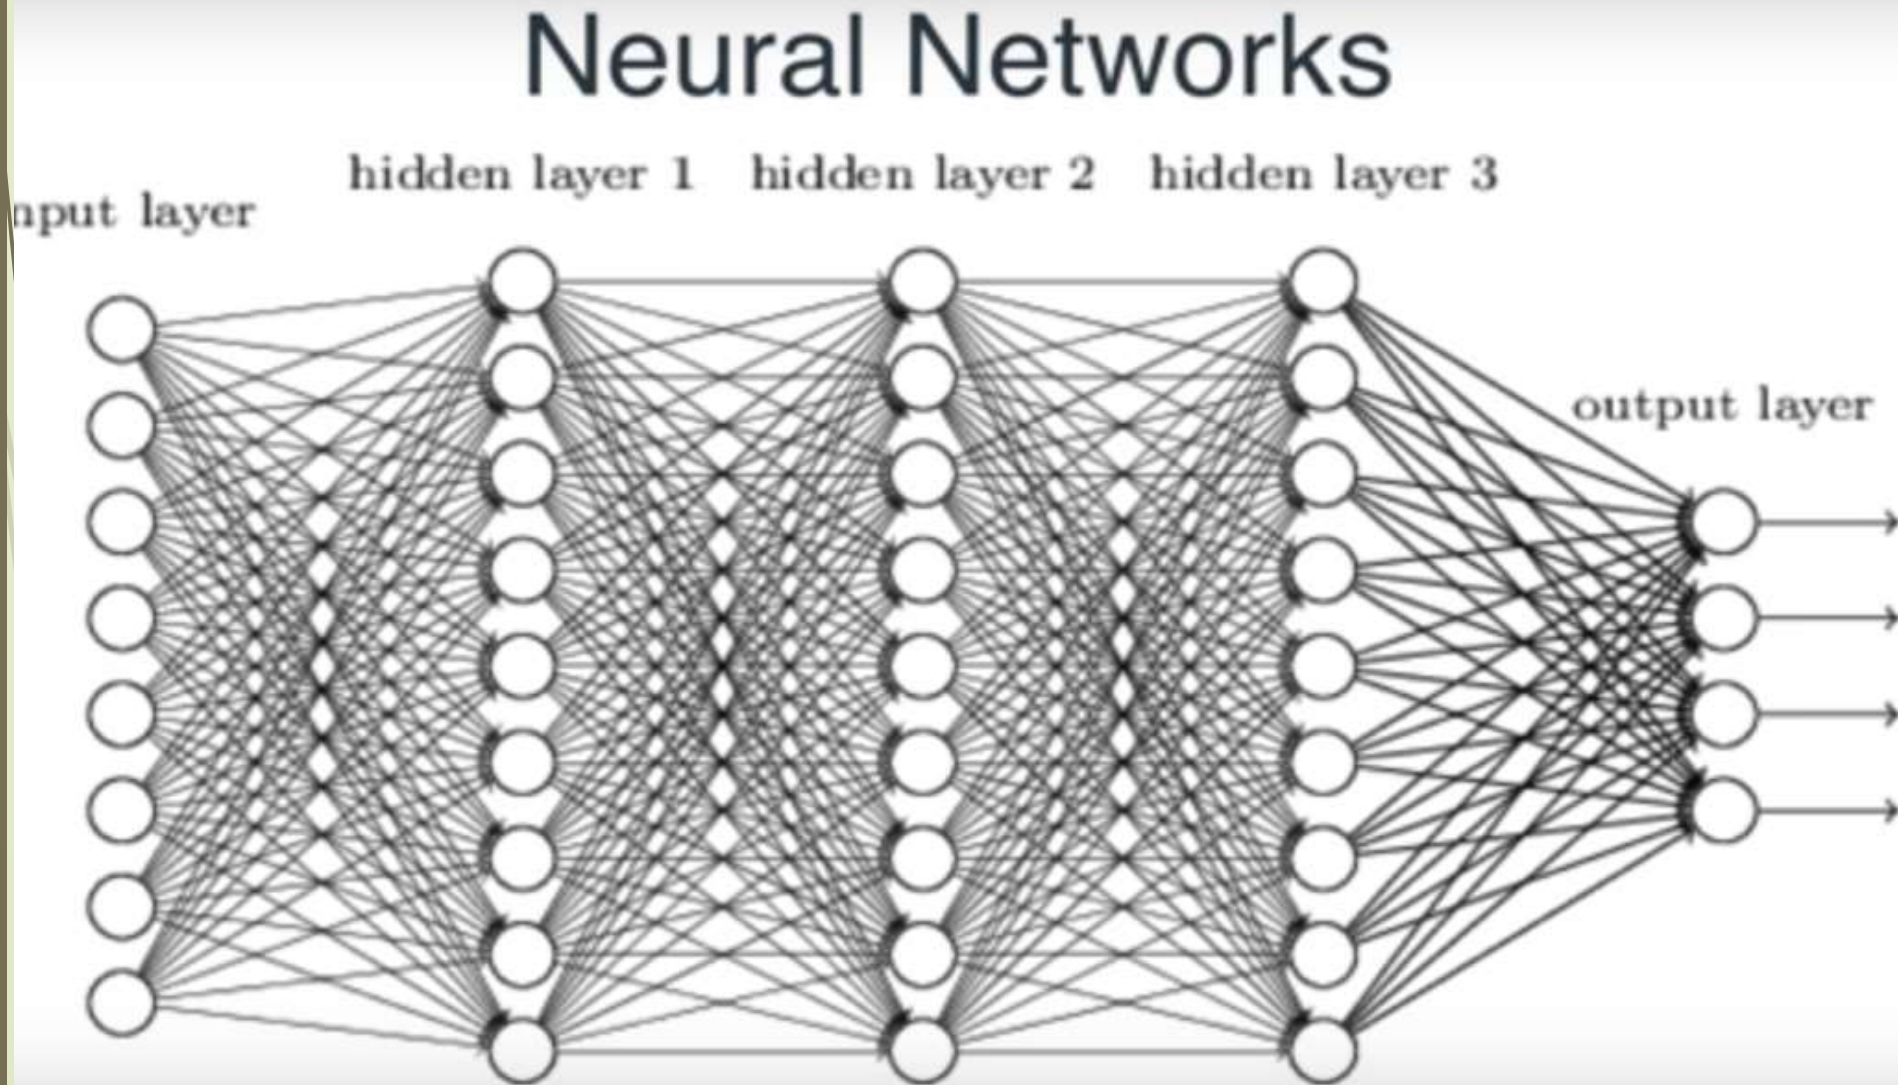
Other TF 2 Namespaces

- ➡ `tf.keras.regularizers` (L1 and L2)
- ➡ `tf.keras.utils` (to_categorical)
- ➡ `tf.keras.preprocessing.text.Tokenizer`

TF 1.x Model APIs

- `tf.core` (“reduced” in TF 2)
- `tf.layers` (deprecated in TF 2)
- `tf.keras` (the primary API in TF 2)
- `tf.estimator.Estimator` (implementation in `tf.keras`)
- `tf.contrib.learn.Estimator` (Deprecated in TF 2)


Neural Network: 3 Hidden Layers



TF 2/Keras and MLPs

- import tensorflow as tf
- ...
- model = `tf.keras.models.Sequential()`
- model.add(tf.keras.layers.Dense(10, input_dim=num_pixels, activation='relu'))
- model.add(tf.keras.layers.Dense(30, activation='relu'))
- model.add(tf.keras.layers.Dense(10, activation='relu'))
- model.add(tf.keras.layers.Dense(num_classes, activation='softmax'))
- model.compile(tf.keras.optimizers.Adam(lr=0.01), loss='categorical_crossentropy', metrics=['accuracy'])

TF 2/Keras and MLPs

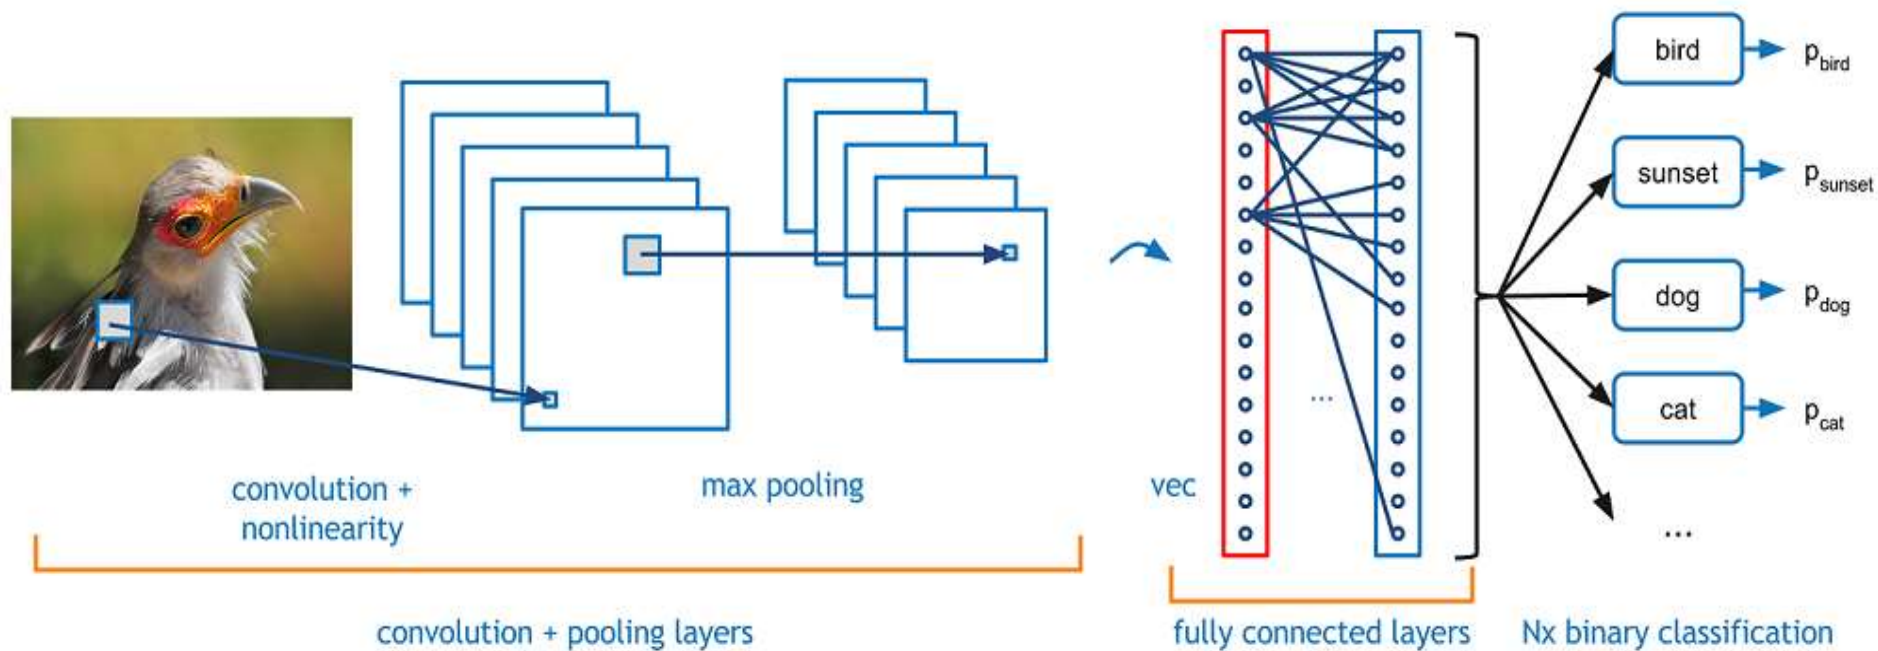


```
import tensorflow as tf


...

model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(10, input_dim=num_pixels, activation='relu'),
    tf.keras.layers.Dense(30, activation='relu'),
    tf.keras.layers.Dense(10, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation='softmax'),
    tf.keras.optimizers.Adam(lr=0.01), loss='categorical_crossentropy',
    metrics=['accuracy'])
]
```

CNNs: Convolution and Pooling



TF 2/Keras and CNNs



```
➤ import tensorflow as tf
➤ ...
➤ model = tf.keras.models.Sequential([
➤     tf.keras.layers.Conv2D(30,(5,5),
➤         input_shape=(28,28,1),activation='relu')),
➤     tf.keras.layers.MaxPooling2D(pool_size=(2, 2)),
➤     tf.keras.layers.Flatten(),
➤     tf.keras.layers.Dense(500, activation='relu'),
➤     tf.keras.layers.Dropout(0.5),
➤     tf.keras.layers.Dense(num_classes, activation='softmax')
➤ ])
```

What are RNNs?

- RNNs = Recurrent Neural Networks
- RNNs capture/retain events in time
- FF networks **do not learn** from the past
- => RNNs **DO learn** from the past
- RNNs contain layers of recurrent neurons

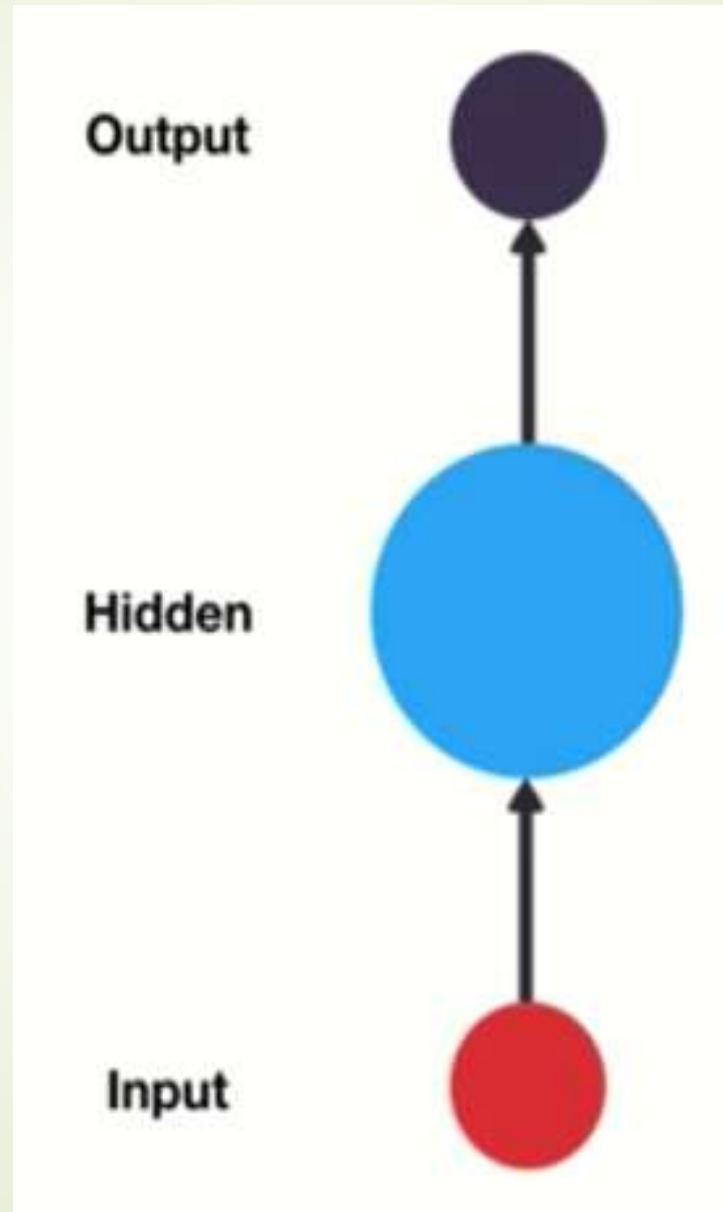
Common Types of RNNs

- 1) Simple RNNs (minimal structure)
- 2) LSTMs (input, forget, output gates)
- 3) Gated Recurrent Unit (GRU)
- NB: RNNs can be used with MNIST data

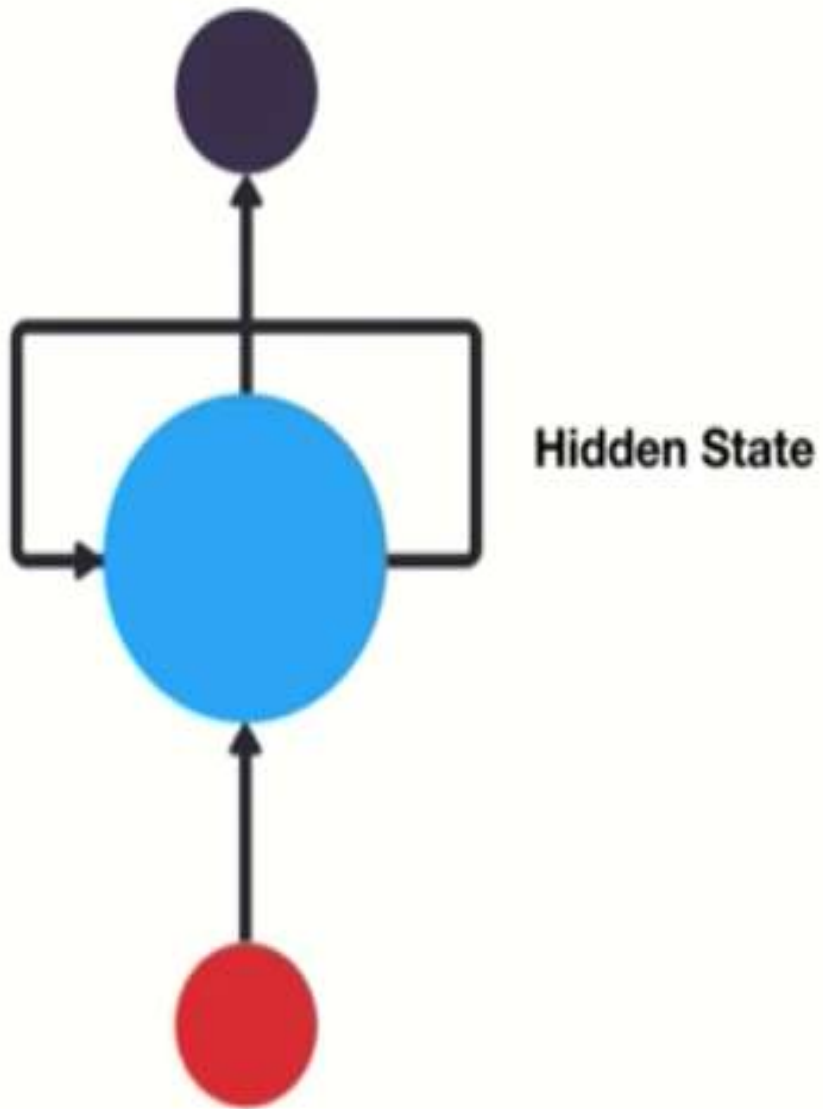
Some Use Cases for RNNs

- ➡ 1) Time Series Data
- ➡ 2) NLP (Natural Language Processing):
Processing documents
“analyzing” Trump’s tweets
- ➡ 3) RNNs and MNIST dataset

Recall: Feed Forward Structure



RNN Cell (left) and FF (right)



RNN “Big Picture” (cs231n)

Recurrent Neural Network

We can process a sequence of vectors \mathbf{x} by applying a recurrence formula at every time step:

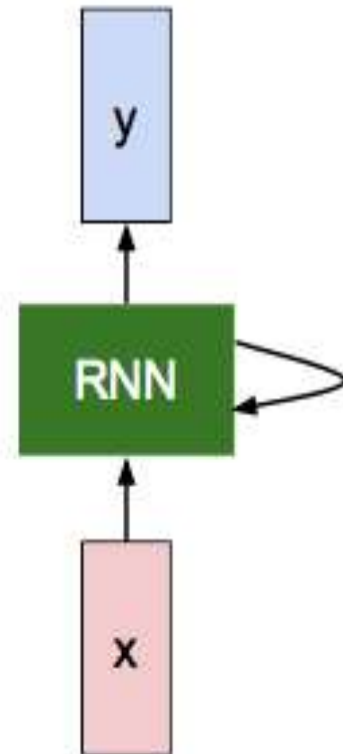
$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

new state

some function with parameters W

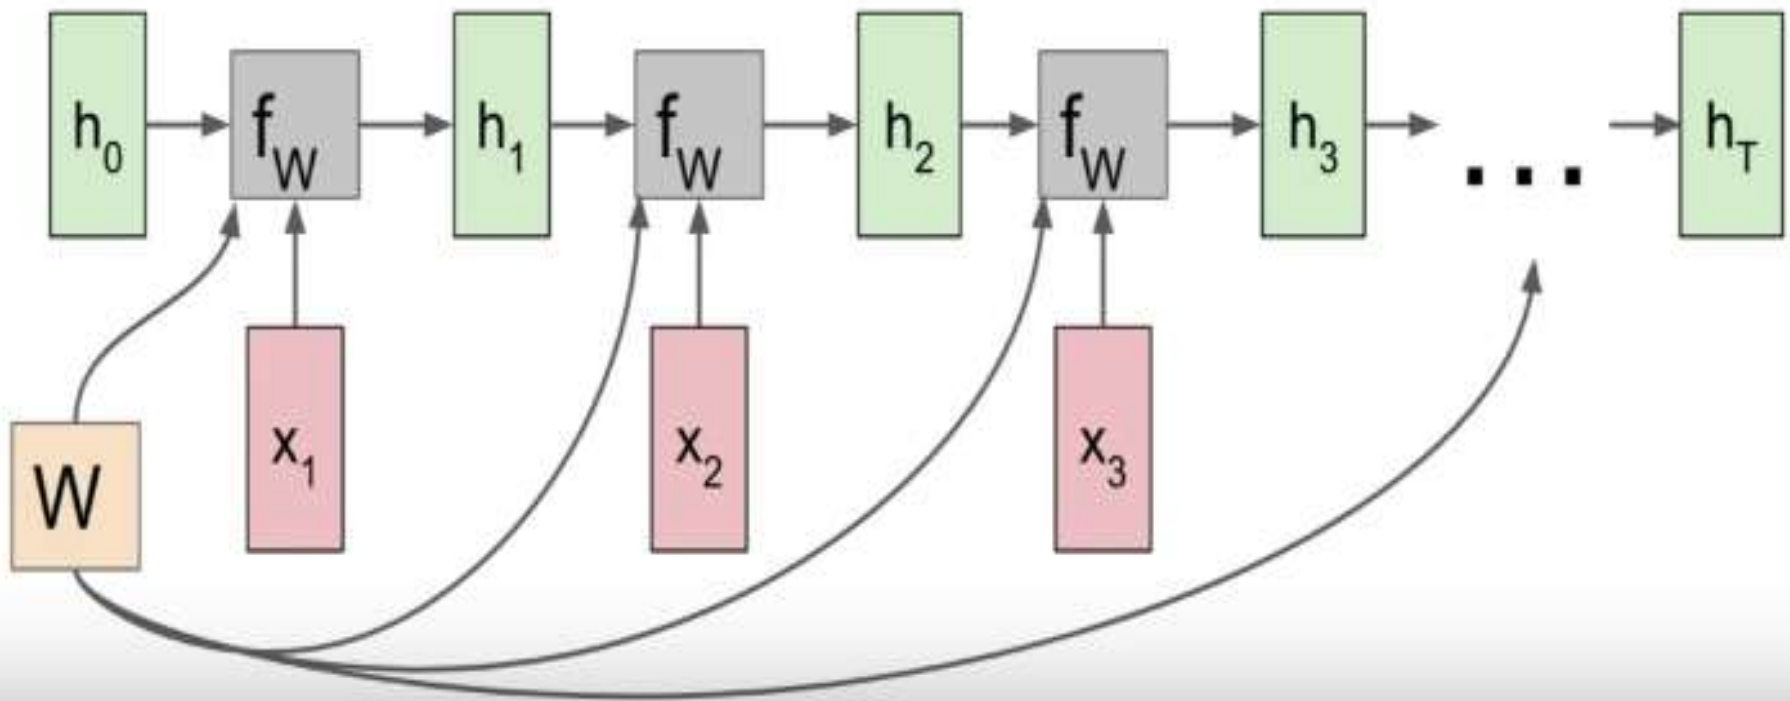
old state

input vector at some time step



RNN Neuron: same W for all steps



Re-use the same weight matrix at every time-step



RNN and MNIST Dataset

- Sample initialization values for RNN:
 - `n_steps = 28` # # of time steps
 - `n_inputs = 28` # number of rows
 - `n_neurons = 150` # 150 neurons in one layer
 - `n_outputs = 10` # 10 digit classes (0->9)
- Remember that:
 - one layer = one memory cell
 - MNIST class count = size of output layer

Problems with RNNs

- 
- ➡ lengthy training time
 - ➡ Unrolled RNN will be very deep
 - ➡ Propagating gradients through many layers
 - ➡ Prone to vanishing gradient
 - ➡ Prone to exploding gradient
- 



What are LSTMs?

- Long short-term memory (LSTM):
a model for the **short-term memory**
Contains more state than a “regular” RNN
which can last for a **long** period of time
can forget and remember things selectively
- Hochreiter/Schmidhuber proposed in 1997
- improved in 2000 by Felix Gers' team
- set accuracy records in multiple apps

Features of LSTMs

- Used in Google speech recognition + Alpha Go
- they avoid the vanishing gradient problem
- Can track 1000s of discrete time steps
- Used by international competition winners

Use Cases for LSTMs

- 
- Connected handwriting recognition
 - Speech recognition
 - Forecasting
 - Anomaly detection
 - Pattern recognition
- 

Advantages of LSTMs

- well-suited to classify/process/predict time series
- More powerful: increased memory (=state)
- Can explicitly add long-term or short-term state
- even with time lags of unknown size/duration between events

The Primary LSTM Gates

- Input, Forget, and Output gates (FIO)
- The main operational block of LSTMs
- an "information filter"
- a multivariate input gate
- some inputs are blocked
- some inputs "go through"
- "remember" necessary information
- FIO gates use the sigmoid activation function
- The cell state uses the tanh activation function

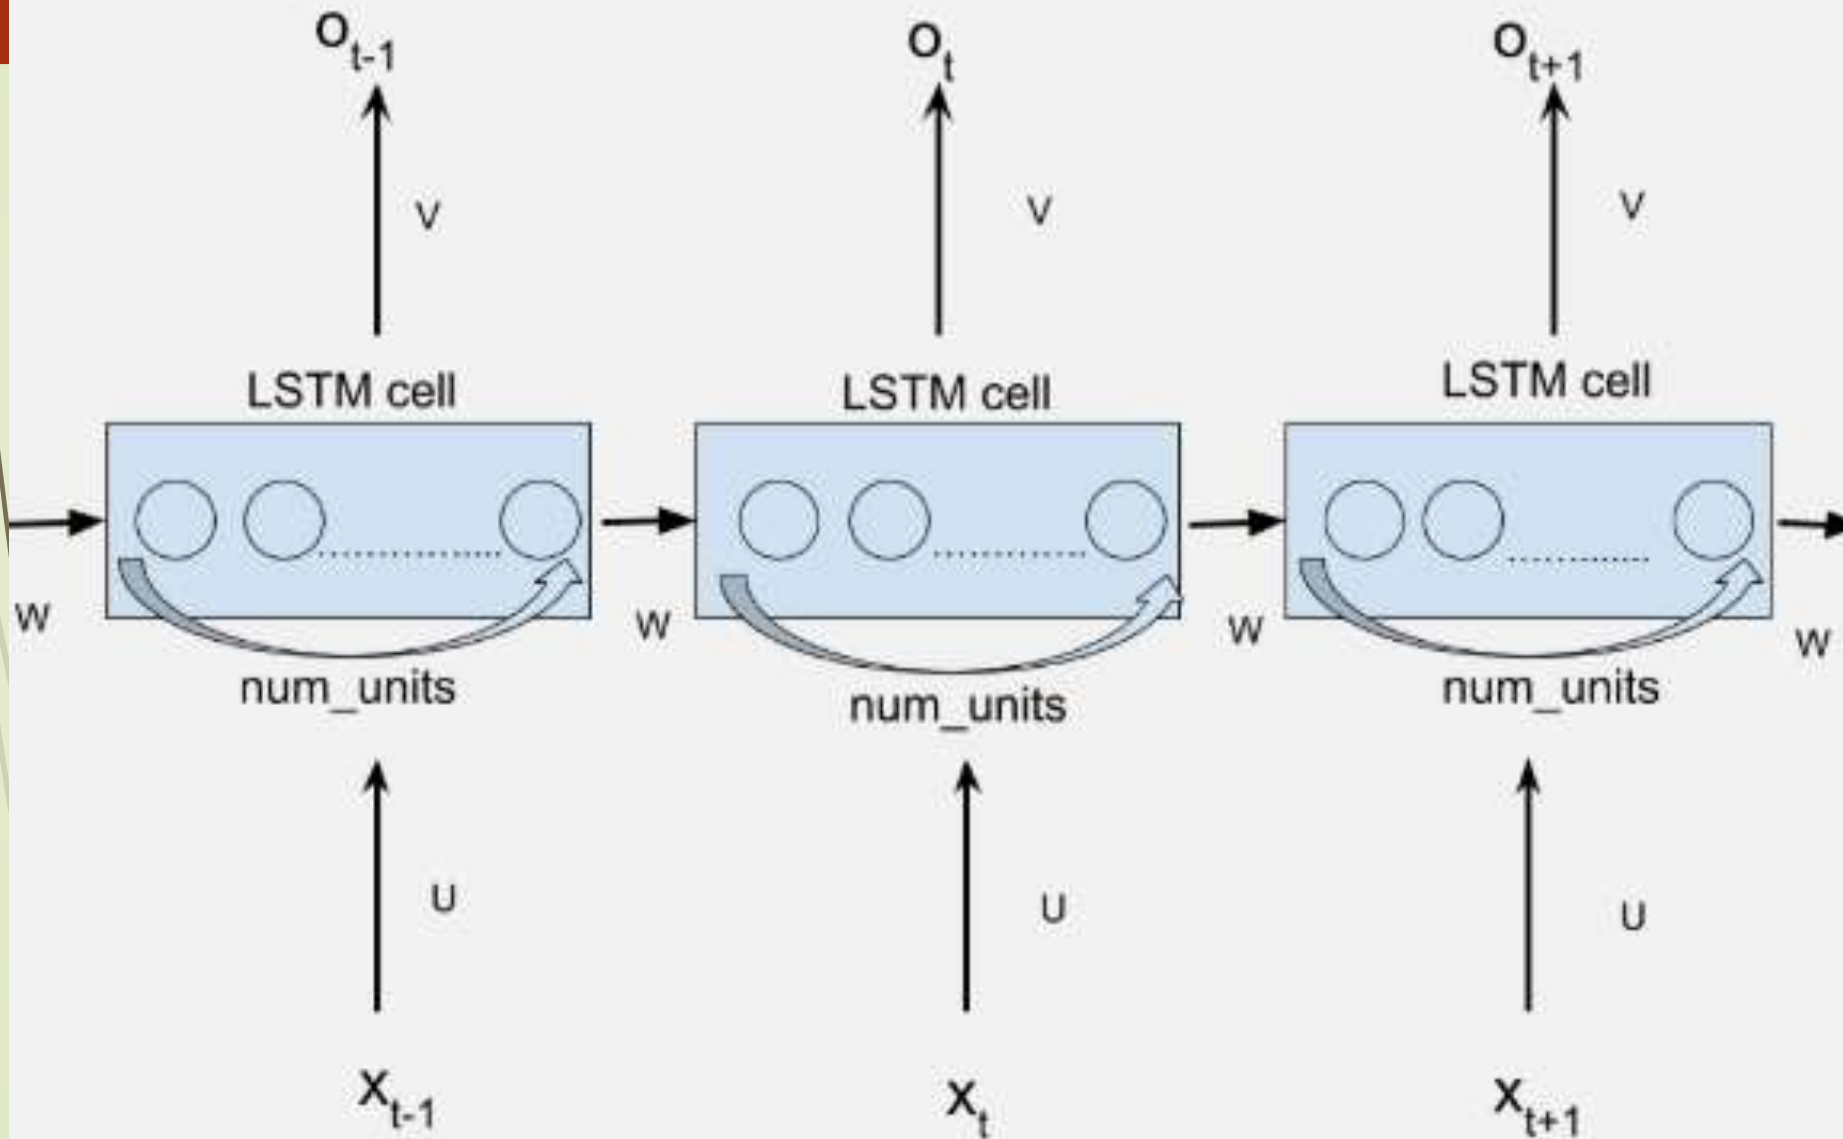
LSTM Gates and Cell State

- 1) the **input** gate controls:
the **extent** to which a new value flows **into** the cell state
- 2) the **forget** gate controls:
the **extent** to which a value **remains** in the cell state
- 3) the **output** gate controls:
The portion of the cell state that's part of the output
- 4) the **cell state** maintains long-term memory
- => LSTMs store more memory than basic RNN cells

Common LSTM Variables

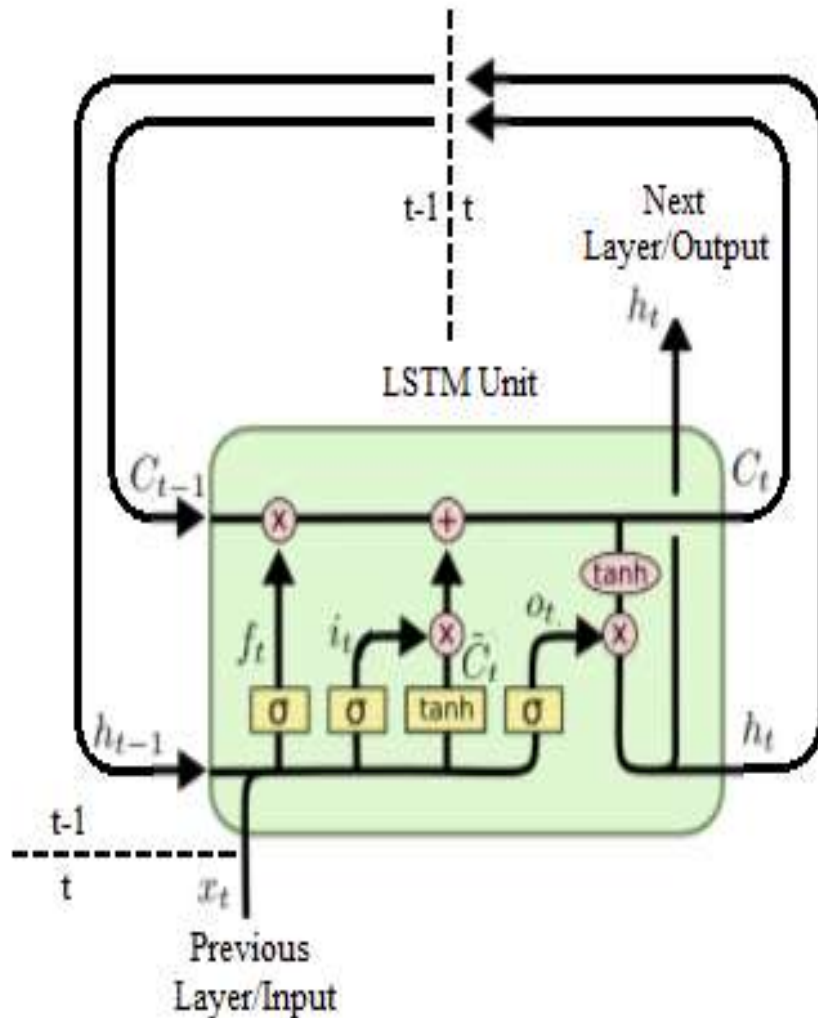
- # LSTM unrolled through 28 time steps (with MNIST):
- `time_steps = 28`
- # number of **hidden** LSTM units:
- `num_units = 128`
- # number of **rows** of 28 pixels:
- `n_inputs = 28`
- # number of **pixels** per row:
- `input_size = 28`
- # mnist has 10 classes (0-9):
- `n_classes = 10`
- # size of input batch:
- `batch_size = 128`

Basic LSTM Cells



Interpreting LSTM cell and num_units

LSTM Formulas



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(C_t)$$

TF 2/Keras and LSTMs

```
import tensorflow as tf
```

```
...
```

```
model = tf.keras.models.Sequential()
```

```
model.add(tf.keras.layers.LSTMCell(6,batch_input_shape=(1,1,1),kernel_initializer='ones',stateful=True))
```

```
model.add(tf.keras.layers.Dense(1))
```

```
...
```

=> LSTM versus LSTMCell:

<https://stackoverflow.com/questions/48187283/whats-the-difference-between-lstm-and-lstmcell>

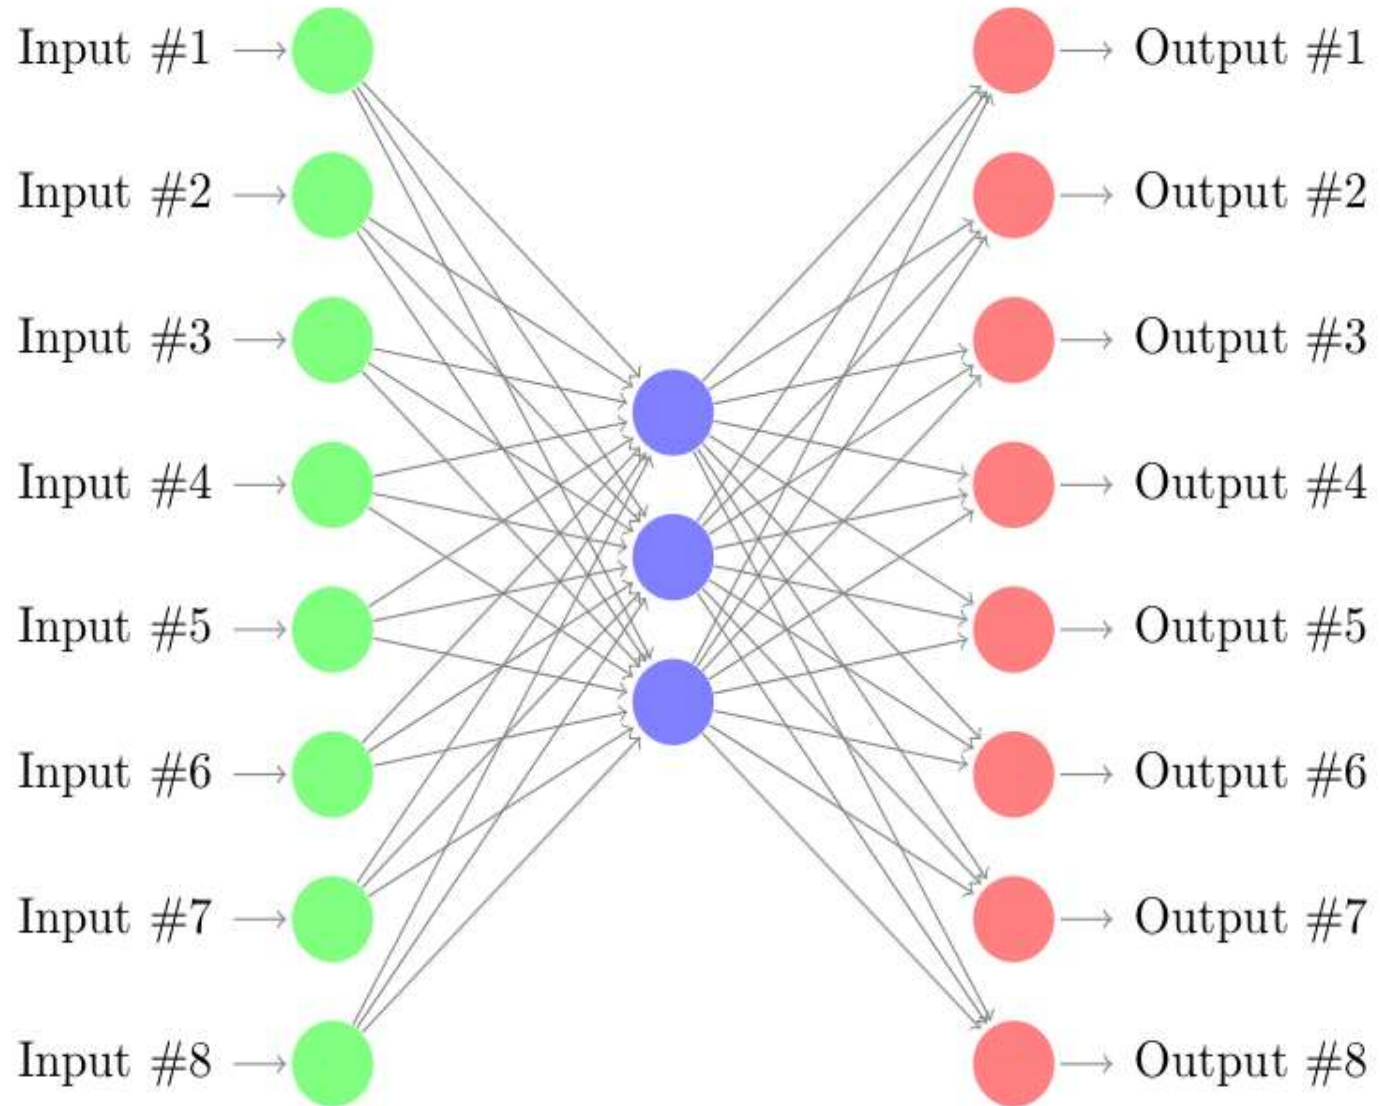
=> How to define a custom LSTM cell:

<https://stackoverflow.com/questions/54231440/define-custom-lstm-cell-in-keras>

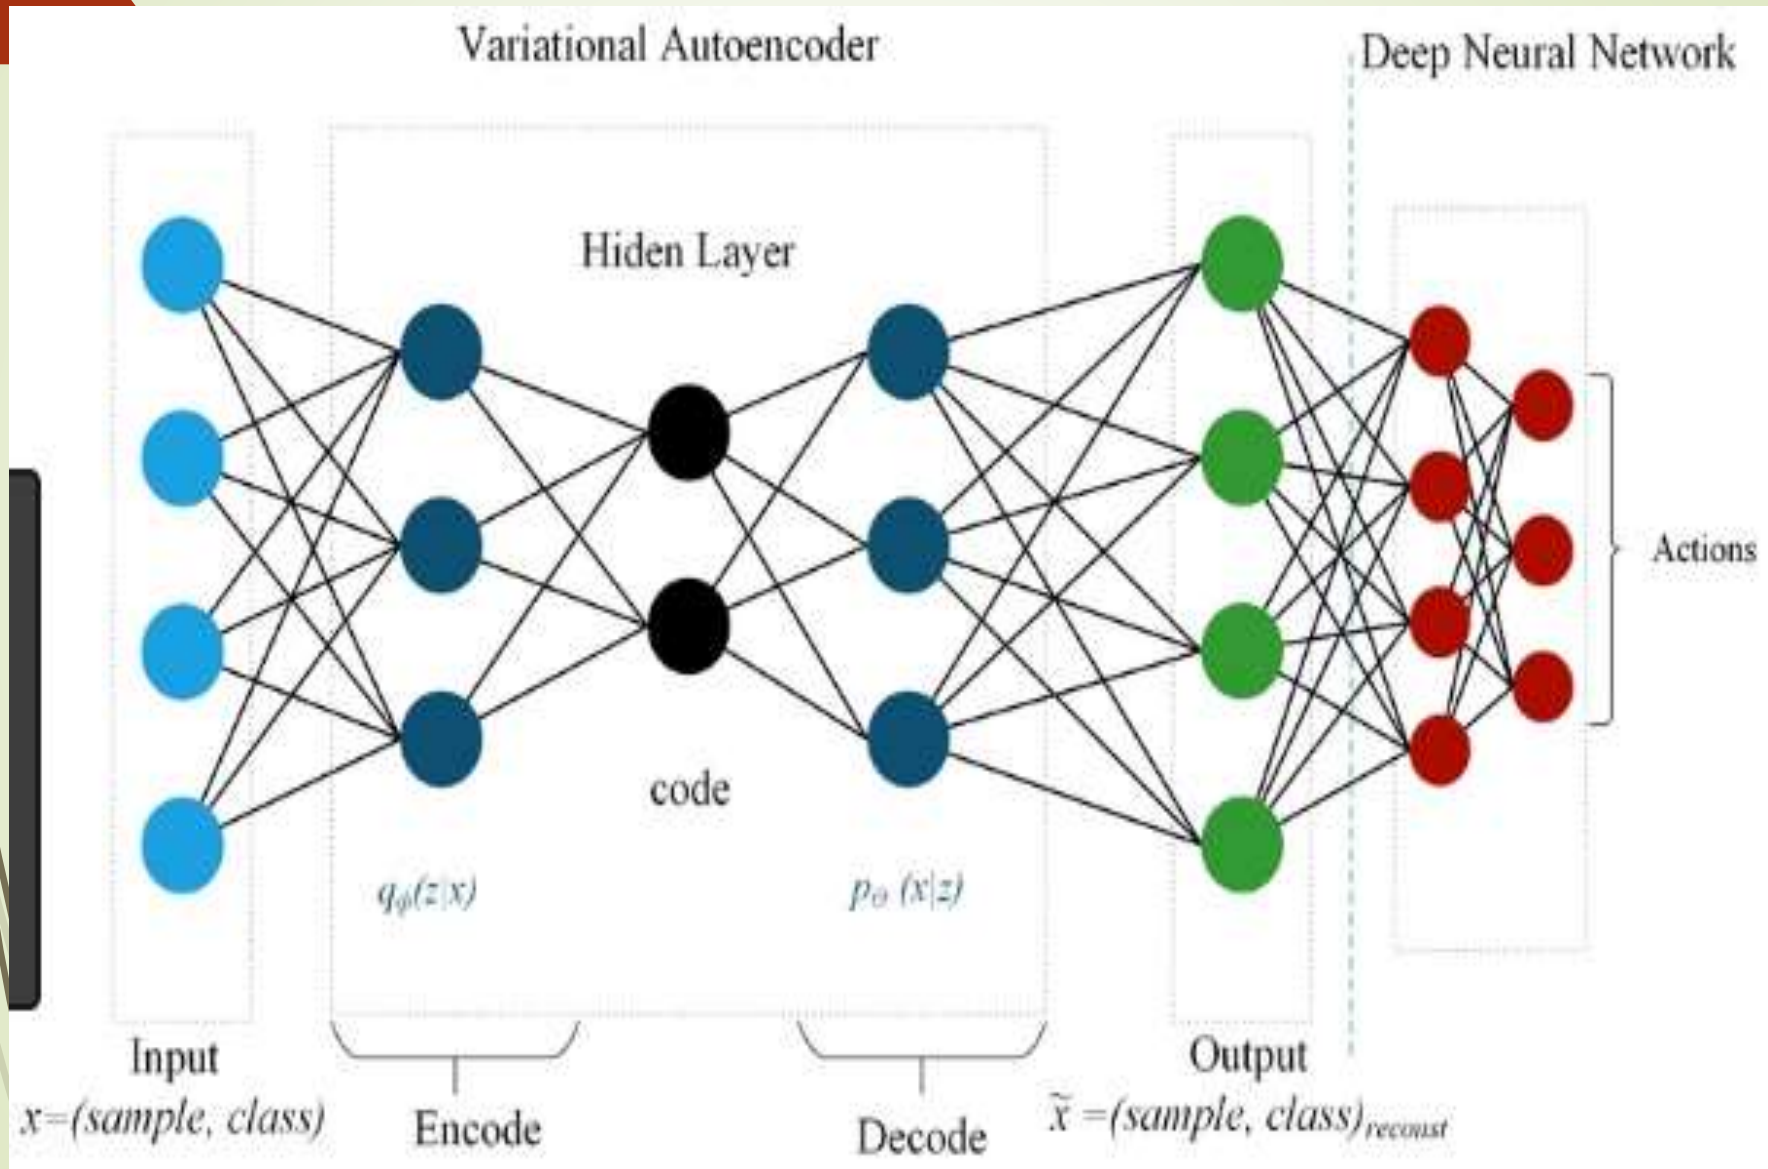
TF 2/Keras & BiDirectional LSTMs

- `import tensorflow as tf`
- `...`
- `model = Sequential()`
- `model.add(Bidirectional(LSTM(10, return_sequences=True), input_shape=(5,10)))`
- `model.add(Bidirectional(LSTM(10)))`
- `model.add(Dense(5))`
- `model.add(Activation('softmax'))`
- `model.compile(loss='categorical_crossentropy', optimizer='rmsprop')`

Autoencoders



Variational Autoencoders (2013)



AEs, VAEs, and GANs

➤ <https://jaan.io/what-is-variational-autoencoder-vae-tutorial/>

➤ TensorFlow 2 and Autoencoders:

<https://gist.github.com/AFAgarap/326af55e36be0529c507f1599f88c06e>

➤ TensorFlow 2 and Variational Autoencoders:

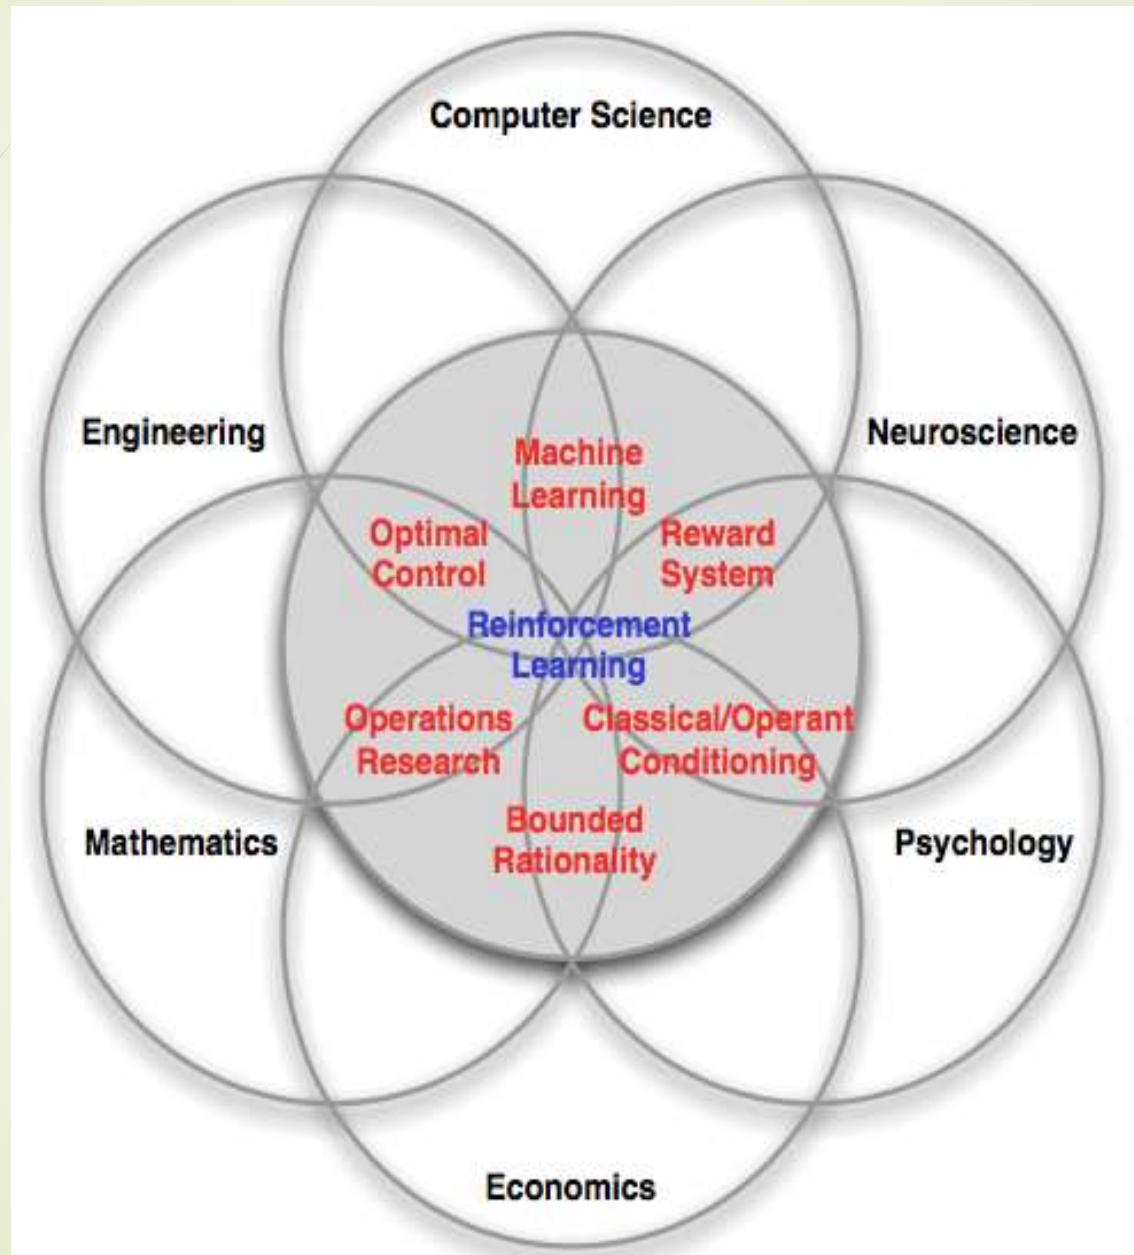
➤ Convolutional Variational Autoencoder:

<https://www.tensorflow.org/alpha/tutorials/generative/cvae>

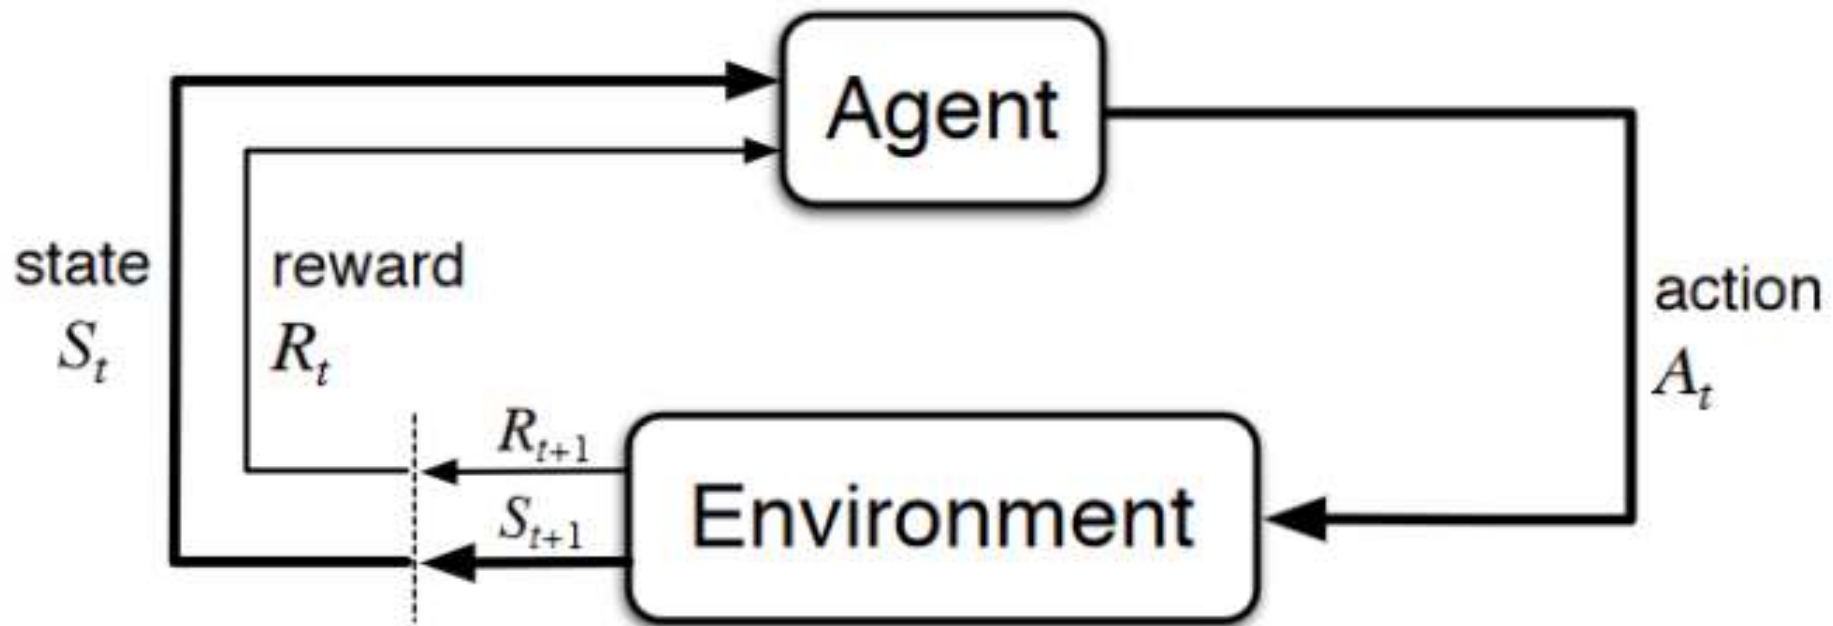
➤ Variational Autoencoders and GANS:

<https://www.youtube.com/watch?v=KFX4apL7a4s>

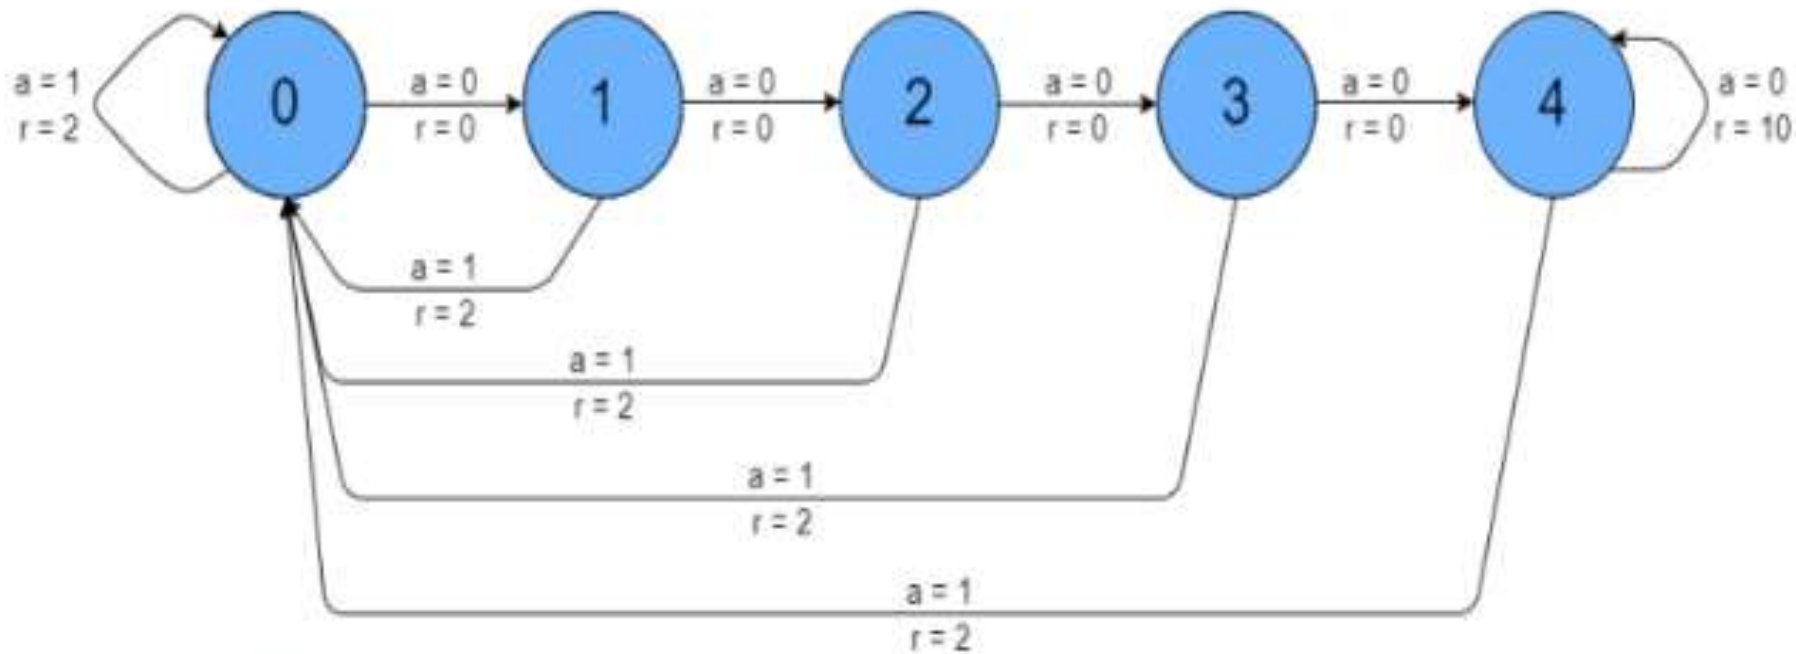
Reinforcement Learning



Reinforcement Learning

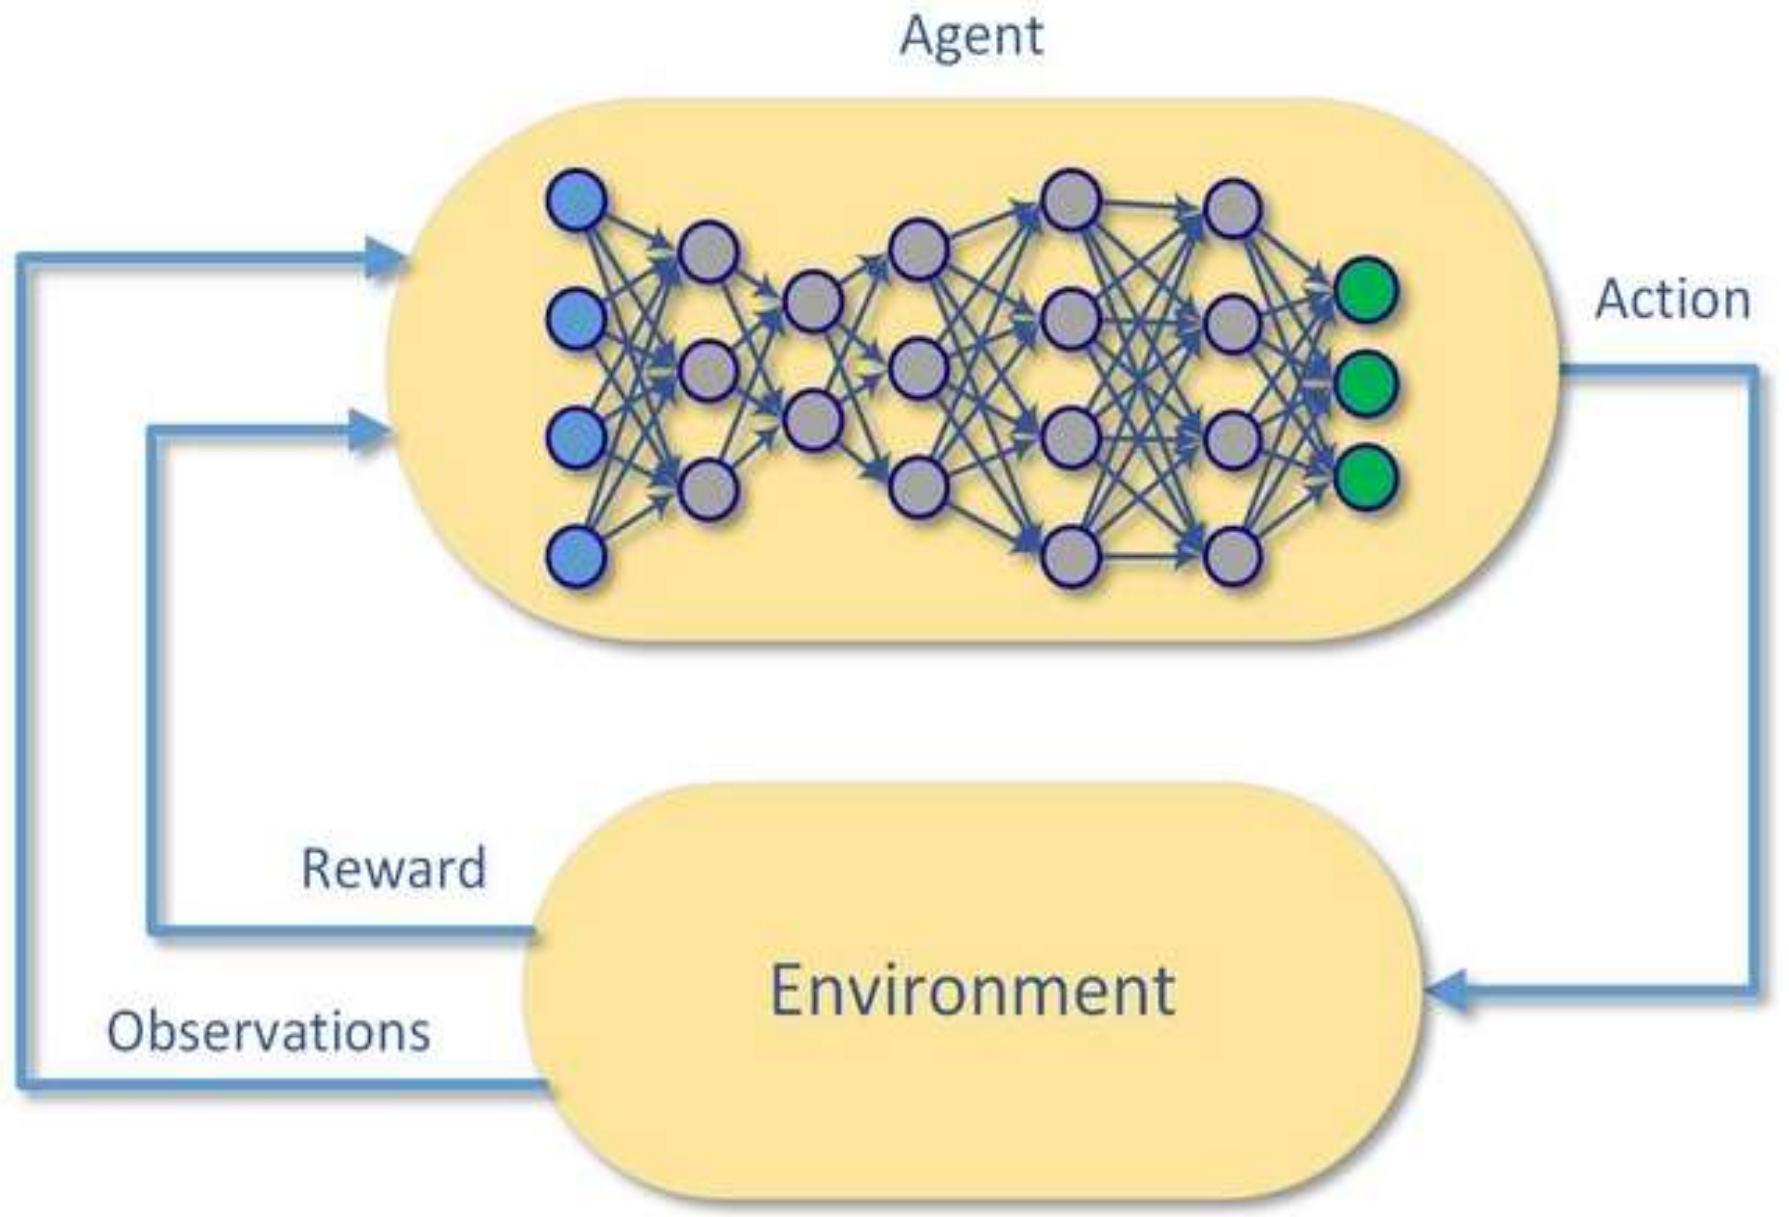


Reinforcement Learning



Open AI Gym's NChain environment

Deep Reinforcement Learning



Reinforcement Learning

- Dopamine on TensorFlow (Research Framework):

<https://github.com/google/dopamine>

- TF-Agents library for RL in TensorFlow:

<https://github.com/tensorflow/agents>

- Keras and Reinforcement Learning:

<https://github.com/keras-rl/keras-rl>

- OpenAI universe, DeepMind Lab, TensorFlow